

Near-beam tagging of the far-forward lithium recoils at ePIC: what the aperture is worth, what the optics is worth, and whether superconducting nanowires help

C. Peng and J. Zhou

Physics Division, Argonne National Laboratory

Writing assisted by Claude (Anthropic)

2026-08-26 · rewritten 2026-08-28 on the per-configuration Yellow Report divergences and the lithium tagging optics, restated 2026-08-29 on the per-configuration transport, corrected 2026-09-02 on the cross-report consistency review, re-sourced 2026-09-06 · figures from `evgen/scripts/nearbeam_aperture_scan.py`, `nearbeam_reach_gain.py`, `nearbeam_sensor_budget.py` and `nearbeam_zid_power.py`; the far-forward geometry from `tools/fullsim` and the eic/epic main branch

Abstract. The two far-forward lithium tags of the programme — the intact ^6Li recoil of the coherent channel and the α spectator of the tagged channels — are rigidity-degenerate with the beam and are separated from it only by their angle at the interaction point, so the coherent tagged fraction is an exponential in the square of the closest approach. This report prices that approach for each beam configuration against the three half-widths that compete to set it: the 10σ envelope of the Yellow Report optics (2.2, 1.8 and 0.92 mrad for ^6Li at 5×41 , 10×100 and 18×275), the aperture of the ePIC Roman-Pot silicon measured by transporting an intact ^6Li through the geometry (2.50, 1.51 and 0.53 mrad, measured on 2026-08-28 in the current layout, the 5×41 figure standing on the ePIC baseline proton lattice and worth $\times 0.64$ in the $Z/A = 0.5$ alternative), and the envelope of a lithium tagging optics in which the horizontal β^* is raised 46–164 $\times$ with the pots following (0.36, 0.19 and 0.12 mrad). At the published optics which of the two binds now depends on the configuration — the silicon at 5×41 , where 2.50 mrad exceeds a 2.20 mrad envelope, and the machine at 10×100 and, by eight orders of magnitude, at 18×275 , where 0.53 mrad is well inside 0.92 — and no closer approach by any technology recovers a tag: at the silicon edge the coherent tag runs 9×10^{-10} to 1×10^{-5} , and once the envelope that binds at 10×100 and 18×275 is applied the best any technology reaches is 7×10^{-8} at 5×41 and 6×10^{-27} at 10×100 , and of the α spectators only the 1.5% off-rigidity slice that the Roman-Pot window catches at any optics together with a 0.1–1.0 point over-rigid branch on the inner side of the bend. Under the tagging optics the silicon becomes the limit and a layer that reaches the envelope is what makes the optics worth having: 25–37% of coherent recoils at 1/6.8–1/12.8 of the luminosity, seven populated $|t|$ bins per configuration through the full reconstruction chain, and a 22–31% α tag. The gain lives in a strip 50–180 μrad wide at the inner edge. Superconducting nanowires do not improve either the approach or the charge identification: their dead edge is worth 0.005 mrad against a gap set by module tiling and by a machine-protection rule, and in charge identification the one bit they deliver is nearly sufficient in principle but the 50% geometric fill of the existing devices cannot reach 95% ^6Li efficiency over four planes, while the $\alpha + d$ breakup that the charge question exists to reject is two hits 3–41 mm apart in sensors ePIC already has — and at the tagging optics, where that background is created, the second hit vetoes 84% of it outright.

1 · The angle is the measurement

The coherent measurement needs the intact ^6Li , whose A/Z is exactly the beam's; the tagged measurements need the α spectator, at 0.99813 of the beam rigidity. Neither is swept out by the dipoles, and the only handle is the transverse angle at the interaction point: the coherent recoil follows $\exp(-B|t|)$ with $B \approx 50 \text{ GeV}^{-2}$ and $|t| = (\theta p)^2$, so the tagged fraction falls as the exponential of the square of the approach angle times the square of the beam momentum. The α 's angular distribution is set by the cluster wave function, with a tail extending to a few hundred μrad at the top energy and a few mrad at the bottom (Figure 1b).

Three half-widths compete to set the approach, and until 2026-08-28 this programme confused two of them by applying one proton-derived, isotropic, energy-independent divergence of 72.7 μrad at every configuration. The **10 σ envelope** is the

operational retraction of the Roman Pots, a rectangle of half-widths $10(\sigma_h, \sigma_v)$ in angle: for ${}^6\text{Li}$ at the Yellow Report high-acceptance optics, 2.2×3.8 , 1.8×1.8 and 0.92×0.92 mrad at 5×41 , 10×100 and 18×275 (Report 3 §4.1) — a machine-protection and halo rule, not a sensor one. The **silicon aperture** is what the sensor package, its shield, its ASIC and its module tiling leave open, measured by transporting an intact ${}^6\text{Li}$ through the ePIC far-forward geometry (`tools/fullsim`): $|\theta_x| \gtrsim 2.50 / 1.51 / 0.53$ mrad against $|\theta_y| \gtrsim (\text{shut}) / 2.12 / 0.92$ mrad, measured on 2026-08-28 in the current layout (§6). And the **tagging-optics envelope** is the 10σ rectangle of the lithium tagging optics of Report 1 §6.1 [13], in which the horizontal β^* alone is raised to the optimum of tagged fraction $\times$ luminosity: 0.36×3.8 , 0.19×1.8 and 0.12×0.92 mrad at $r_h = \beta_x^* / \beta_{x,\text{HA}} = 46.5$, 164.1 and 89.3 and $L/L_{\text{HA}} = 1/6.8$, $1/12.8$ and $1/9.5$ (`farforward.tagging_optics_point`). In the aperture scan of §2 the vertical half-width is the larger of the measured aperture and the envelope, per configuration; in the chain of §2.1 the pots at the silicon use the measured rectangle and the pots that follow use the tagging optics' 10σ rectangle in both planes, as Report 1 §6.1 prices it.

2 • What a closer approach is worth




Figure 1 — the price of every aperture. Tagged fraction of the coherent intact ${}^6\text{Li}$ (a, $\exp(-B|t|)$, $B = 50 \text{ GeV}^{-2}$, whose 40–60 band is the model uncertainty) and of the ${}^6\text{Li}$ α spectator routed through the far-forward systems (b, cluster model, $\beta = 0.30 \text{ GeV}$ — the short-range parameter of the two-parameter cluster momentum density, whose 0.20–0.40 band is the model uncertainty) against the horizontal half-width of the near-beam cutout in angle, per configuration, with the vertical half-width fixed as in §1. Open circles: the silicon aperture of the current geometry, measured 2026-08-28; squares: the 10σ envelope of the Yellow Report high-acceptance optics; triangles: the envelope of the tagging optics. The curves are technology-independent and price any closer approach by any means; the markers are where the three constraints sit — except the two 10×100 coherent ones, 2.0×10^{-19} at the silicon and 6.2×10^{-27} at the Yellow Report envelope, which fall below panel (a)'s 10^{-18} floor and are read from Table 1 instead. `nearbeam_aperture_scan.py`.

CONFIGURATION (${}^6\text{Li}$ GEV/U)	HALF-WIDTH	5 × 41 (40.8)	10 × 100 (99.5)	18 × 275 (137.5)
coherent recoil, tagged fraction	silicon, 2026-08-28 (2.50 / 1.51 / 0.53 mrad)	9.4×10^{-10}	2.0×10^{-19}	1.2×10^{-5}
	YR high-acceptance envelope (2.20 / 1.80 / 0.92)	7.2×10^{-8}	6.2×10^{-27}	7.1×10^{-14}
	tagging-optics envelope (0.36 / 0.19 / 0.12)	0.37	0.25	0.33
${}^6\text{Li}$ α spectator, tagged (any far-forward system)	silicon	0.017	0.016	0.029
	YR high-acceptance envelope	0.018	0.016	0.025
	tagging-optics envelope	0.31	0.22	0.29

Table 1 — the three half-widths priced (Figure 1's markers; `nearbeam_aperture_scan.py`, re-run 2026-08-29 on the per-configuration transport; the silicon half-width it prices was measured 2026-08-28). The α row is the tag routed through

the far-forward systems (`farforward.acceptance_summary`), the same routing as Report 3 Table 6 but at this scan's vertical half-width, $\max(\text{measured aperture}, 10\sigma_v)$ — 6.49 mrad at 5×41 , where the measured plane is shut, and 2.12 mrad at 10×100 , against Table 6's 3.80 and 1.80 mrad envelopes, the two coinciding to 0.3% at 18×275 , 0.92 against 0.9169 mrad. That difference is worth almost nothing, and what it is worth runs the other way: the near-beam band is tagged only outside the rectangle, so a wider vertical admits marginally *fewer*, and on a common sample the two verticals differ by 0.007 points at 5×41 , 0.003 at 10×100 and nothing at 18×275 , against 0.001 points at the tagging optics. What separates these markers from Table 6's in the last digit is the Monte-Carlo draw and not the aperture: this scan draws 2×10^5 spectators per marker from a fresh `default_rng(7)` against Table 6's 4×10^5 from a stream earlier channels have already consumed, and at that sample size the seed-to-seed spread is 0.04–0.06 points on the Yellow Report row and 0.07–0.11 on the tagging row, seed 7 the lowest of twenty at all three configurations at once, and the two rows are one quantity with Table 6's — 0.0184 / 0.0168 / 0.0259 and 0.3147 / 0.2231 / 0.2910 at 8×10^6 spectators. The 1.5% off-rigidity slice below $R = 0.95$ is tagged at any half-width, the over-rigid branch adds 0.1 to 1.0 points on the inner side of the bend, and the rest is the $R \approx 1$ tail outside the rectangle. The tagging optics is paid for in luminosity, 1/6.8, 1/12.8 and 1/9.5 of the high-acceptance value, so its rows are to be read as yield per unit high-acceptance luminosity $\times 0.147$, $\times 0.078$ and $\times 0.106$. Two statements follow from the table and neither depends on a sensor. At the Yellow Report optics the silicon is now outside the envelope at 5×41 , 2.50 against 2.20 mrad, and inside it at the other two, so moving the silicon closer buys a factor 77 at 5×41 and nothing at 10×100 or 18×275 — the coherent tag stays below 10^{-7} , and of order 10^{-5} at the silicon edge at 18×275 where the silicon is now the more permissive of the two by eight orders of magnitude (both at $B = 50 \text{ GeV}^{-2}$; over the 40–60 band the factor runs 33 to 180 and the two numbers behind those thresholds move by $\times 20$ and $\times 7$ either way), and the α tag stays at the 1.5% of its off-rigidity slice plus the over-rigid branch. At the tagging optics the envelope shrinks by 6–9 $\times$ and the silicon, now at 0.53–2.50 mrad, becomes the only limit: pots that follow the envelope tag 25–37% of coherent recoils and 22–31% of α spectators, pots that stay where they are tag the off-rigidity slice alone. The layer and the optics are multiplicative levers, and the rest of this report prices the layer given the optics. The top energy is no longer excluded. Under the single 0.727 mrad envelope of the earlier analysis the coherent channel at 18×275 required $p_T > 0.60 \text{ GeV}$ and was dead at 10^{-9} for any aperture; at the tagging optics the envelope is 0.12 mrad, $p_T > 0.10 \text{ GeV}$, and the top configuration is the most productive of the three at equal luminosity. The earlier ordering of the configurations was an artefact of the divergence used.

2.1 · Through the chain

Acceptance is not the measurement. Figure 2 runs the full coherent chain of Report 2 §4.5 — the Roman-Pot emulation with the anisotropic divergence and the rectangular cutout, the two azimuths $\alpha = \varphi_e - \varphi_S$ and $\beta = \varphi_t - \varphi_S$, the spin-state-sorted two-dimensional harmonic fit — at the tagging optics of each configuration, once with the pots fixed at the measured silicon aperture and once with the pots following the envelope, at the tagging optics' luminosity ($10 \text{ fb}^{-1}/u$ per year $\times L/L_{\text{HA}}$) and $P_{zz} = 0.60$.




Figure 2 — the chain at the tagging optics. (a) δa_t per $|t|$ bin at one year with the pots following the envelope, for the three configurations; the silicon-aperture runs contribute no point to either panel at any configuration, the 18×275 one because its single populated $|t|$ bin falls under the script's minimum-count guard rather than because it is empty (Table 2). (b) The recovered a_t against the injected deformation term. The abscissa of both panels is each bin's t_{ref} , the rate-weighted mean true $|t|$ at which a_t is quoted (Report 2 money plot 6R(d)), and not the reconstructed edges of Table 2: the seven bins sit at 0.023–0.092 GeV^2 at 5×41 , 0.025–0.083 at 10×100 and 0.023–0.114 at 18×275 . `nearbeam_reach_gain.py`, 2×10^6 Monte-Carlo recoils per run.

CONFIGURATION	POTS	CUTOUT [MRAD]	ACCEPTANCE	$N_{\text{TAG}} / \text{YR}$	$ t $ BINS	ΔA_T PER BIN	A_E (INJECTED 0.0100)
5 × 41	silicon	2.50 × 6.49	0	0	0 of 7	—	—
5 × 41	follow	0.36 × 3.80	0.36	2.3×10^6	7 of 7	0.0063, 0.0052, 0.0051, 0.0037, 0.0050, 0.0089, 0.0190	$0.0125 \pm 0.0022 \dots$ -0.0034 ± 0.0141
10 × 100	silicon	1.51 × 2.12	0	0	0 of 7	—	—
10 × 100	follow	0.19 × 1.80	0.25	2.4×10^6	7 of 7	0.0130, 0.0072, 0.0063, 0.0040, 0.0048, 0.0077, 0.0148	$0.0100 \pm 0.0023 \dots$ -0.0114 ± 0.0103
18 × 275	silicon	0.53 × 0.92	1.5×10^{-5}	268	0 of 7	—	—
18 × 275	follow	0.12 × 0.92	0.32	6.0×10^6	7 of 7	0.0039, 0.0031, 0.0032, 0.0023, 0.0033, 0.0064, 0.0153	$0.0115 \pm 0.0013 \dots$ 0.0273 ± 0.0115

Table 2 — what the aperture costs the measurement, at the tagging optics. The seven $|t|$ bins are 0.017–0.028, 0.028–0.039, 0.039–0.05, 0.05–0.08, 0.08–0.12, 0.12–0.17 and 0.17–0.25 GeV^2 , the window adopted 2026-08-28 (Report 2 §7); a_t is quoted at each bin's t_{ref} , the rate-weighted mean true $|t|$, and wherever that falls below the lowest point of its digitized anchor at 0.05 GeV^2 the injected a_t is the linear deformation shape extrapolated — four of the seven bins and not three, the resolution putting the t_{ref} of the 0.05–0.08 bin at 0.046 GeV^2 , 0.0498 at 18×275 (Report 2 §4.5). The injected a_t and its errors are that section's modulation coefficient $c_2 = 2a_2$ per unit P_{ZZ} ; the shape-term significance of §7 is on the a_2 footing instead. a_e is the exotic-glue coefficient, recovered consistent with its injected value in every bin, nowhere more than 2.1 standard errors away. The recovered a_t falls below the injected curve in the highest bin at 5×41 and 10×100 — 0.1196 ± 0.0190 against 0.1835, -3.4σ , and 0.1167 ± 0.0148 against 0.1666, -3.4σ — while at 18×275 it is consistent at 0.2298 ± 0.0153 against 0.2282, $+0.1\sigma$ — a fluctuation of that one draw rather than the absence of the defect, whose twenty-experiment ratio mean in the same bin is 0.2060 against 0.228 injected, against the likelihood's 0.2289 (Report 2 Table 5). That is the measured low-count bias of the bin-wise ratio estimator — the $O(1/\nu_b)$ Jensen offset of inverting $R = \sigma_p^2 T / (1 +$

$u + \bar{P}T$) at 45 counts per (α, β) cell, which the acceptance-profiled likelihood `reco.harmonic_likelihood_fit_2d` removes outright (Report 2 §4.3 and Table 5, plans/08 A12), not a property of the spectrum at the cutout's edge; every other bin agrees to within 1.7σ , the largest pulls -1.65σ at 18×275 in 0.05–0.08 and -1.51σ at 5×41 in 0.12–0.17. The analytic yields of Report 1 §6.1, 2.4, 2.4 and 6.1×10^6 per year, are reproduced to 3% by the Monte-Carlo rectangle. The one silicon row that is not identically dead buys nothing, and the 1.5×10^{-5} it reads is the divergence-smeared Monte-Carlo reading of the rectangle Table 1 prices analytically at 1.2×10^{-5} , resting on 29 accepted recoils of the 2×10^6 drawn and so carrying $\pm 19\%$: 268 tagged recoils a year at 18×275 put 231 into a single $|t|$ bin, which is a count and not a measurement — under two per (α, β) cell of the design — and since 2026-08-28 the script's `MIN_TAGGED_PER_BIN = 1000` drops it rather than report the $a_t = -1.56 \pm 2.22$ it would fit at a t_{ref} outside the linear model's validity.

With the pots fixed the measurement does not exist at any configuration: the silicon half-aperture stands 6.9, 7.9 and 4.5 envelope widths out horizontally against 1.7, 1.2 and 1.0 vertically, and the acceptance is zero to the precision of 2×10^6 recoils at the two lower configurations and 1.5×10^{-5} at the top, $\pm 19\%$ on the 29 recoils that fraction rests on — where it is not identically zero only because the horizontal inner edge, re-measured at 0.53 mrad against the 1.03 mrad of the earlier reading, is the shallowest of the three in transverse momentum, 0.44 GeV/c against 0.61 and 0.90, or 4.4 standard deviations of the $\exp(-B|t|)$ recoil rather than 6.1 and 9.0. The vertical half-height is not the binding cut at any configuration: it is an upper limit, and 0.92 mrad at 18×275 is 7.6 standard deviations out. With the pots following, every configuration returns seven $|t|$ bins with $\delta a_t = 0.0023\text{--}0.0130$ in the four below 0.08 GeV², alongside which Report 1 §6.1 prices the counting floor of its best (x, Q^2) super-bin alone at 5σ floors of 1.7, 2.3 and 1.6% per unit P_{zz} per year and 3.0, 5.5 and 2.6 years to 5σ on a 1% exotic-glue term; the seven bins together measure that flat term to $\delta a_e = 0.00121, 0.00111$ and 0.00074 per year (Report 1 §6.3). The earlier conclusion that the coherent channel is a low-energy programme — Report 2's before its 2026-08-28 rewrite, and withdrawn with it — was conditioned on the single 0.727 mrad envelope; at the tagging optics the top configuration is the best covered, as Report 2 §5.2 now has it.

The same optics decides what the *tagged* observables can see, and there the currency is not acceptance but reach. The α -tagged tensor asymmetry $A_{zz}^{\text{tag}}(k)$ of the embedded deuteron — the S/D interference of the α -d relative motion, and the first tagged spin observable projected for any $A > 2$ — is routed through the same envelope at 10×99.5 with 4×10^5 events of the S + D tagged sampler (`polligen.tagged`, through `money_tagged_azz.py`), whose α tag is 2.5% at the Yellow Report optics against the 1.5% Roman-Pot slice of Table 1's S-wave row — the D wave and nothing else, reconciled in Report 3 Table 6 — and 25% at the tagging optics, which at $L/L_{\text{HA}} = 1/12.8$ costs 19% in tagged events per year (0.0247 against 0.0199) — and 10×99.5 is the only configuration at which it costs anything. What that 19% buys is where in k those events sit. At the Yellow Report optics the only surviving α spectators are the off-rigidity slice that falls below $R = 0.95$, which is the high- k tail of the density: the median accepted spectator momentum is 0.32 GeV/c and **not one accepted event lies below $k = 0.15$ GeV/c** at any of the three configurations (medians 0.348, 0.323, 0.334 GeV/c). At the tagging optics the near-beam tail reopens and the median falls to 0.153–0.177 GeV/c with 36–48% of the sample below 0.15 — against an unfolded density whose median is 0.093 GeV/c with 74% below 0.15. On the same sampler $\text{acc} \times L/L_{\text{HA}}$ runs $0.0284 \rightarrow 0.0503$ at 5×41 and $0.0265 \rightarrow 0.0331$ at 18×275 , so at the other two configurations the reach arrives with 1.8 and 1.2 times the tagged rate rather than a cost. Whether the $O(1)$ asymmetry Cosyn and Weiss predict at $p_s \approx 300$ MeV/c [14] survives the acceptance is therefore answered "yes, and at the published optics *only* there": the tagging optics is what turns a one-point measurement into a curve. It also changes what the curve means, because the accepted sample is sculpted in the polar angle θ_k of the spectator and the two optics sculpt it oppositely — the off-rigidity slice is longitudinal ($|\cos \theta_k| = 0.80$) and the near-beam tail transverse (0.39) — so the folded asymmetry reads +0.49 and -0.07 at $k \approx 0.33$ GeV/c where the analytic $\theta_k = 90^\circ$ curve says -0.48 . Comparing either against the 90° curve reads the interference at the wrong angle; the acceptance-weighted prediction reproduces both, to $|\Delta A_{zz}| \leq 0.063$ at the Yellow Report optics and ≤ 0.085 at the tagging optics in every populated k bin — within 1.1σ and 1.6σ at the 4×10^5 events of that sample — and, at 8×10^6 , where the errors have fallen to 0.003–0.03, to ≤ 0.025 and ≤ 0.022 , every bin then within 1.6σ and 2.2σ of the bin-averaged prediction.

For ${}^7\text{Li}$ the same trade runs backwards, and that is the sharpest programme statement in this report. The ${}^7\text{Li}$ α is off rigidity at $R = 0.856$, in the *middle* of the Roman-Pot momentum window, so it is tagged by the window and never has to clear the

near-beam envelope at all: its acceptance is 0.97 at the Yellow Report optics and 0.99 at the tagging optics, flat to three points across the whole 0.05–3 mrad aperture axis — in angle. At 5×41 that routing is optimistic: the per-energy insertion holds the station-1 silicon band off to $|y| \geq 18$ mm, so an α at $y \approx 0$ reaches silicon only at station 2, on a module edge, in 57% of events in the full simulation of the current geometry (Report 3 Table 6, `tools/fullsim`); at 10×100 and 18×275 the insertion is nearer and the angle-only figure stands. A tagging optics re-optimised on the ${}^7\text{Li}$ beam ($\beta_x^* \times 62, 220$ and 102) therefore buys ${}^7\text{Li} \times 1.02$ in acceptance — at 5×41 on the angle-only routing, the per-band insertion having been measured at the published optics only — for $\times 1/7.9$ to $\times 1/14.8$ in luminosity — a factor 7.7–14.5 *net loss*, which at equal running time multiplies every ${}^7\text{Li}$ error bar by 2.8, 3.8 and 3.2 at 5×41 , 10×100 and 18×275 ; in the ${}^6\text{Li}$ tagging optics of §1, the only new one the programme asks for, ${}^7\text{Li}$ pays $\times 1/6.8$ to $\times 1/12.8$ instead and its bars grow by 2.6, 3.5 and 3.0. The in-situ alignment polarimeter and the tagged polarized-EMC companion are free at the published optics and cost a factor 7.7–14.5 at the ${}^7\text{Li}$ -optimised one. That companion is also where the γ^2 target-mass term of the longitudinal asymmetry is largest in the programme: with the exact finite- γ $A_{||}$ the kernel has carried by default since 2026-08-29 (`InclusiveKernel(target_mass=True)`; `target_mass=False` is the massless limit), the σ -weighted shift of the tagged-triton overlay reaches 2.09% in the top x bin at 10×99.5 and 5.04% at 5×40.8 , against $\leq 0.6\%$ and $\leq 2.5\%$ at the money-plot sweet spots of those same two configurations (Report 1 §2), because the $-\gamma^2 g_2/F_1$ inside A_1 dominates as $g_2^{\text{WW}} \rightarrow -g_1$ near $x \approx 0.5$ (`target_mass_bound.py`). The inversion of the programme's own strategic finding — that ${}^7\text{Li}$'s tag is the easy one and ${}^6\text{Li}$'s the hard one — is thus optics-independent in both directions: ${}^6\text{Li}$ needs the de-squeeze and ${}^7\text{Li}$ must not have it. The two isotopes want different machine optics, which makes them different runs rather than one fill plan, and that belongs alongside the two questions for C-AD in §7.

A trap the earlier analysis exposed and this one removes. The reconstructed-level coherent money plot carried `cut_scale_x = 2.5` from a pre-measurement belief that the pots surround a wide horizontal slot; with the aperture measured, keeping it imposed a 25σ retraction no machine rule asks for, binding before the geometry did. Every script of this report holds the envelope at 10σ per axis, applies the rectangle to each fragment's azimuth, and lets the measured geometry be the only other aperture.

3 • Where the gain lives, and why the dead edge is not it

The exponential piles the gain against the inner edge. Moving the horizontal half-width from the measured silicon to the tagging envelope, half of the recovered coherent yield lies within 176, 65 and 50 μrad of the envelope at the three configurations, 90% within 489, 184 and 141 μrad , and 99% within 823, 315 and 239 μrad (`nearbeam_sensor_budget.py`); for the α spectator at 137.5 GeV/u the same widths are 77, 270 and 624 μrad . A near-beam insert is a narrow strip, not a plane: at the measured $R_{12} = 29.97$ m, taken to survive the de-squeeze (§6), the 50% strip at 18×275 is 1.5 mm wide, and a 3 mm strip along 20 mm of vertical extent on both sides of the beam is 120 mm² — 960 channels of 1×1 mm² microwire tiles, or 133 000 nanowire devices of 30×30 μm^2 . The gap the layer has to cover is set by the measured aperture rather than by the retired 72.7 μrad envelope, and at 18×275 it is 12.4 mm — 0.53 mrad in to 0.12 — so 496 mm² and 3968 microwire channels.

Near-beam recoil tagging is on record as the first of the four EIC applications of superconducting nanowires listed in Yellow Report §14.5.3 — "a Roman pot detector in the forward region about 35 meters or more from the interaction point to tag low momentum transfer recoiling ions" [1] — and the FY22 generic-detector R&D proposal states the mechanism: "The so-called 10σ rule-of-thumb for Roman pots prevents the beam from damaging the detectors by limiting its proximity to the beam. Radiation hard nanowire detectors will be able to move closer to the beam, thus extending the acceptance reach at lower t " [2]. Neither half of that mechanism survives the arithmetic. The 10σ rule protects the beam and keeps the sensor out of the halo — the ePIC geometry file itself calls its per-energy insertions "the ten-sigma cuts for the Roman pots, translated to the physical layout we currently have" [3] — and a more radiation-hard sensor does not relax it; the LHC ran the CT-PPS pots at 25σ at 6.5 TeV. And the dead-edge argument, while it describes a real sensor property — a nanowire's active area reaches within a few hundred nanometres of the substrate edge [1], sharper than the 10 μm telescope used to look at it [8] — is worth ≈ 150 μm against the 100–200 μm of the slim-edge planar and 3D silicon already deployed in Roman Pots: 0.005 mrad at the measured $R_{12} = 29.97$ m, 4% of the tagging envelope at the top energy and

1.2% of the 0.41 mrad gap between the silicon and that envelope. What does survive is the granularity half of the argument, and it belongs to ePIC rather than to a detector-R&D group: the insertion is quantised by the module, a 16 mm block against a tagging envelope of 3.52 mm at 18×275 , and the fix is a staircase retraction, narrower modules, or the x-moving layer ePIC is already designing — "one 'x' moving layer and one 'y' moving layer in each station" [5]. That layer, following the envelope of the tagging optics, is the entire content of Table 2.

4 • Charge identification

An intact ${}^6\text{Li}$, an α and a d from breakup share rigidity and velocity, so only dE/dx separates them, as z^2 (9 : 4 : 1); no EIC document addresses Z identification in the pots, and the only documented concept is a z^2 Cherenkov behind the IR-8 secondary-focus pots. Open question #19 of the programme asks whether the pots can tell an intact ${}^6\text{Li}$ from its breakup.

A nanowire will not do it by pulse height, because the device latches: the amplitude is the diverted bias current, and both Fermilab/JPL beam papers find the amplitude distributions of 120 GeV hadrons and muons, 8 GeV pions and showering electrons indistinguishable [7,8]. **It would do it by firing threshold**. In the normal-core hot-spot regime a deposit Q makes a normal core of radius $r_s \propto \sqrt{Q}$ and the wire switches when the core spans it against the bias-suppressed superconducting width,

$$r_s = \sqrt{\frac{Q}{e\pi c\rho(T_c - T_0)}}, \quad \frac{I_{th}}{I_c} = 1 - \frac{2r_s}{w}$$

a relation three groups reach from three directions — Argonne's Eqs. 1–2 for 120 GeV protons [4], Renema's photon form, and the ion literature's $I_{th}/I_c = 1 - (zeV)^{1/2}C/w$ [9]. Because $dE/dx \propto z^2$ at fixed velocity, and a ${}^6\text{Li}$ at 137.5 GeV/u has $\beta\gamma = 148$ against 128 for Argonne's calibration proton, r_s scales linearly in z with no velocity correction: 134 nm (p, d), 268 nm (α), 402 nm (${}^6\text{Li}$) on Argonne's anchor.




Figure 3 — Z by firing threshold. I_{th}/I_c against wire width for the three species from $r_s \propto z$ anchored on Argonne's 134 nm; open circles are their measured thresholds, digitised from Fig. 7 of [4]; dotted verticals mark $w = 2r_s$, below which the wire fires at any bias and carries no charge information. Inverting the measured points individually gives $r_s = 102\text{--}120$ nm (mean 113) against the quoted 134 nm; every design number is a ratio between species and the offset moves nothing (`evgen/tests/test_nearbeam.py`). `nearbeam_sensor_budget.py` .

WIRE WIDTH	I_{TH}/I_C , P/D	A	${}^6\text{Li}$	SEPARATED TURN-ONS
250 nm	0.00	0.00	0.00	none — Argonne's MIP optimum
400 nm	0.33	0.00	0.00	none
800 nm	0.67	0.33	0.00	α from d
1000 nm	0.73	0.46	0.20	all three — the existing microwire
1500 nm	0.82	0.64	0.46	all three

Table 3 — turn-on bias per species and wire width. Argonne's efficiency optimum, $w \approx 250$ nm, is $2r_s$ for a proton: maximal MIP efficiency and zero charge information. Z identification needs deliberately wide wires — the 1 μm microwire Caltech/JPL/Fermilab already fabricate [7] — where two bias points at 0.33 and 0.60 I_c tag Z by the firing pattern alone, both below the 0.80 I_c dark-count wall and below the photon threshold, which is what a device looking at a 300 K beam pipe needs anyway.

Three caveats weigh on the model. $r_s = 134$ nm is the extrapolated zero-crossing of a four-point fit whose authors write that "the physical validity of this simple model is a question of future work" [4]; the volume in which Q is counted differs by 5000 between Argonne's two papers — the 20 eV film deposit in 2024, the 0.1 MeV substrate deposit in 2026 [5] — which leaves $r_s \propto z$ intact but not the absolute radii; and nobody has varied Z at fixed velocity on one of these devices, Argonne's ${}^{241}\text{Am}$ α differing from their proton through $1/\beta^2$ and every ion demonstration entering charge as $E = zeV$ rather than as z^2 in dE/dx .

4.1 · Against the incumbent: bits versus fill factor

EICROC digitises, per channel, an 8-bit charge from a 40 MHz successive-approximation ADC alongside its 25 ps time-of-arrival, behind an AC-LGAD with a 30 μm active thickness at 500 μm pitch [11]. The right figure of merit is not a separation in σ — a Landau upper tail falls as 1/λ, nothing like a Gaussian — but a fake rate at matched efficiency, with the Landau sampled (`nearbeam_zid_power.py` : four planes, 95% ⁶Li efficiency, 1.5 × 10⁶ events, 6.7 × 10⁻⁷ statistical floor):

READOUT	A FAKE RATE	⁶ LI EFFICIENCY
8-bit charge, per-plane likelihood ratio — the Neyman-Pearson optimum	2.3×10^{-5}	0.950
one bit per plane, majority-of-k, 99% efficient planes	3.1×10^{-5}	0.950
truncated mean — the standard analogue dE/dx estimator	2.7×10^{-3}	0.950
plain sum of the four planes	5.3×10^{-2}	0.950

Table 4 — fake rate at matched efficiency. One bit costs a factor 1.4 against the optimum; a truncated mean costs a factor 86 against one bit, because a Landau has no mean and a single δ ray drags it. The species are far apart (MPV 31.7 against 75.2 keV) and the discriminating power comes from requiring coincidence across planes, not from precision within one.




Figure 4 — fill factor, not bits. The coincidence that makes one bit sufficient needs every plane to record the track. A silicon pixel plane does so ≈ 99% of the time; a superconducting wire comb only over its geometric fill — 25% (SMSPD [7]), 40% [8], 50% (the Argonne EIC-targeted device [4]). At 50% a four-plane array tops out at 93.7% ⁶Li efficiency and cannot reach 95% at any working point. `nearbeam_zid_power.py` .

The nanowire therefore loses on a fabrication number rather than an information-theoretic one, which is the one item on this list the Argonne group could change. The unknown that outranks the comparison is the incumbent's own: EICROC's 8-bit full-scale range is unpublished. A ⁶Li deposits 75 keV = 3.3 fC unamplified, ≈ 67 fC at LGAD gain 20 (α: 28 fC), while EICROC0 is documented as "sensitive to low charge (2 fC)" and characterised at 12.5 fC; if the pot front end is MIP-ranged the ⁶Li clips and α and ⁶Li pile up at full scale, and LGAD gain suppression for highly ionising particles, measured at up to ≈ 2.5×, compresses them toward each other besides (Report 3 Table 9, item 5).

4.2 · The limit on any dE/dx scheme, and the handle that is free

The δ -ray upper tail is the same for every technology and nearly independent of sample thickness; thickening a plane does not help, independent planes do, because an α needs four upward fluctuations rather than one — which is why the four-plane rates of Table 4 sit at 10^{-5} while any single plane is at the percent level. The rates are conservative: the Landau is sampled untruncated, but a δ ray above ≈ 50 keV escapes a $30\text{ }\mu\text{m}$ layer.

The background #19 exists to reject, ${}^6\text{Li} \rightarrow \alpha + \text{d}$, is separable without dE/dx at all. Both fragments sit at beam rigidity — the α at 0.99813 and the deuteron at 1.00452, mass-to-charge ratios rather than mass-number ratios — but the breakup's relative momentum is transverse and unboosted, and the α at $4p_u$ and the d at $2p_u$ take opposite kicks of the same k_T . Their lab angles are therefore in the inverse ratio of their masses, $\theta_d \rightarrow 1.987\theta_\alpha$ as $k \rightarrow 0$, at azimuth $\varphi_\alpha + \pi$; away from that limit the ratio is not fixed — it runs from 0.7 to 5 over the sampled k — but the deuteron is the wider fragment in all but 3×10^{-5} of breakups, and the α is the binding one, so the single-hit topology is overwhelmingly a deuteron alone (6–7.5% of breakups at the Yellow Report optics) rather than an α alone (1.5–1.7%). Sampling both fragments from one relative momentum ($\kappa = 60.7\text{ MeV/c}$; the 3^+ resonance's 42.2 MeV/c is a different production mechanism), the topology at the pots is

CONFIGURATION	MEDIAN SEPARATION	16-84%	IN 500 MM PIXELS	BOTH FRAGMENTS RECORDED YR / TAGGING	A ALONE FAKES A COHERENT TAG YR / TAGGING	PARTNER D THERE TO VETO IT YR / TAGGING	RECORDED PAIR INSIDE ONE PIXEL, TAGGING
18 × 275	10.9 mm	3.2–26.9	22	0.0002 / 0.22	0.0007 / 0.27	0.25 / 0.84	4.2×10^{-3}
10 × 100	10.7 mm	3.4–26.4	21	0.0000 / 0.17	0.0001 / 0.21	0.02 / 0.82	2.2×10^{-3}
5 × 41	17.3 mm	6.3–41.2	35	0.0002 / 0.25	0.0020 / 0.30	0.12 / 0.84	1.5×10^{-4}

Table 5 — the $\alpha + \text{d}$ topology at the pots (`nearbeam_two_hit.py` , both fragments boosted from one relative momentum: 4×10^5 breakups per configuration for the quantiles and 1.2×10^7 for every fraction, because at the Yellow Report optics the α fakes a tag once in 5×10^2 to 10^4 breakups and a small sample conditions on a handful of events — the YR veto column rests on 8.7×10^3 , 1.4×10^3 and 2.4×10^4 fakes at 18×275 , 10×100 and 5×41 and is 0.245 ± 0.005 , 0.017 ± 0.004 and 0.120 ± 0.002 ; the tagging column on $\approx 3 \times 10^6$ each). "YR" is the high-acceptance setting; the high-divergence one is the same beam at 5×41 and a wider envelope at the other two, where the α fakes a tag in only 4×10^{-6} – 4×10^{-5} of breakups and no partner deuteron was recorded at all in 1.2×10^7 . An intact ${}^6\text{Li}$ is one hit; the breakup is two, tens of pixels apart, in sensors that exist. The millimetre is a lever plus a dispersion, $x = R_{12}\theta\cos\varphi + D(R - 1)$, with (R_{12} , R_{34} , D) measured per configuration — 19.24 / 4.56 / 0.311, 21.25 / 3.35 / 0.287 and 29.97 / 2.93 / 0.292 m — and the cluster density's short-range scale $\beta = 0.20$ – 0.40 GeV moving the medians by -12% to $+9\%$ (the veto fractions by less than 0.03). R_{34} is a factor 4 to 10 under R_{12} , which is why every separation here fell by up to 55% against the single-lever numbers. Its 5×41 entry took a second geometry to obtain, the 29.6 mm insertion having shut the vertical plane and left the first scan nothing to regress: the ladder was repeated through a copy of the compact file whose four Roman-pot section offsets are zero, which slides the silicon onto the beam axis and leaves every field alone, and gives $R_{34} = 4.56\text{ m}$ with $R_{12} = 19.18\text{ m}$ against the 19.24 measured through the real insertion as its control. Until then this row carried the fallback $R_{34} = R_{12}$ and read 25.8 mm. The dispersive term is not a correction but half the answer, and it is why a merge is rare rather than impossible: the two fragments take opposite k_z , so their rigidities move apart rather than together, and $D(R_\alpha - R_d)$ has a median of 9.5, 8.8 and 8.9 mm at 5×41 , 10×100 and 18×275 . In the angular lever alone a recorded pair stays $\approx 3\text{ min}(R_{12}\text{ env}_x, R_{34}\text{ env}_y)$ apart — the two levers differing by 4.2× at 5×41 , 6.3× at 10×100 and 10.2× at 18×275 , and the minimum taken over the axes a recorded α is seen to use rather than over both, since at 5×41 no recorded α reaches the 3.80 mrad vertical envelope at either optics — measured minima 20.0, 11.7 and 8.1 mm at the tagging optics against a scale of 21, 12 and 8.1 mm, within 10% either way because $\theta_d/\theta_\alpha = 1.987$ only as $k \rightarrow 0$ — and nothing merges; the dispersion is free to cancel that, and does in the last column. The two lower separation rows supersede the 73.4 and 30.1

mm computed at the retired rigidity-scaled energies, all three the 38.4 / 18.5 / 15.1 mm of the single-lever era (5×41 first, as above), and the 5×41 row the 25.8 mm it read while its vertical lever was the fallback $R_{34} = R_{12}$ (Appendix A).

The second-to-last column is the result. Conditioned on the α landing in the near-beam tail — the only way this channel fakes an intact ${}^6\text{Li}$ — the partner deuteron is on a pot in 84% of events at the tagging optics, so hit multiplicity alone rejects five sixths of the background the coherent measurement's whole charge-identification discussion exists for, with sensors ePIC already has and no dE/dx at all. At the Yellow Report optics the same veto is worth 0.02–0.25, but there the α fakes a tag only once in 5×10^2 to 10^4 breakups, so there is almost nothing to veto: the optics that creates the background supplies the veto with it. One caveat is structural rather than statistical. The 84% assumes a pot that accepts out to $\theta = 5$ mrad, which is 96, 106 and 150 mm at the measured per-configuration R_{12} and is an acceptance-table convention, not a sensor. The outer edge has since been measured, and it is nearer: the debris-free primary track stops at 2.85, 3.85 and 4.00 mrad (`tools/fullsim`), past which the ion has struck the pipe. At those edges the veto is 0.79, 0.82 and 0.84, so the convention costs nothing at the two upper configurations and five points at 5×41 . Bringing the edge further in is what breaks it: against 2.0, 1.0 and 0.5 mrad — 38, 19 and 10 mm at 5×41 , 42, 21 and 11 mm at 10×100 and 60, 30 and 15 mm at 18×275 , the last of these one 16 mm module — the veto falls to 0.68 / 0.22 / 0.00 at 5×41 , 0.82 / 0.69 / 0.21 at 10×100 and 0.84 / 0.81 / 0.56 at 18×275 . At 5×41 the median separation given an α tag is 35 mm against 19 and 17 mm at the other two, so how much of this veto is real is a question about where the pot stations sit and how far they extend — the geometry of §6, measured since (question B1).

5 • What stands in the way of a cryogenic sensor in a pot

OBSTACLE	THE NUMBER	STATUS
cold-stage power	NbN needs 2.8–3 K, a WSi microwire 0.8 K, and an SNSPD cold stage carries a few hundred μW to a few mW; a 4 K pulse tube delivers 1.5 W at 4.2 K and nothing at its 2.8 K floor, on a 190 kg compressor drawing 10.7 kW	hard
beam-induced RF heating	a ferrite-shielded TOTEM pot dissipates ≈ 40 W at the LHC; scaling $Q^2 M / \sigma_z^{3/2}$ to the EIC gives $\approx 16 / 8 / 0.8$ W at highest- E_{CM} / highest-L / 41 GeV — the first two are 11 $\times$ and 5 $\times$ a pulse tube's whole 4.2 K lift, and the last, half of it at 4.2 K, must be removed at the 2.8–3 K the film needs, where the same cooler is at or near its no-load floor	hard
in-vacuum precedent	none for a biased superconducting detector inside an accelerator's beam vacuum; the CryoBLM and AD SQUID precedents sit outside it, and the AD beam is $\approx 10^7$ antiprotons against 8×10^{13} protons	open
vibration	cold-head displacements 7–9 μm , 10 nm with a decoupling structure that is heavy and stiff — the opposite of a pot on motorised rails	mitigable
radiation hardness	unmeasured: the FY22 deliverable was "an upper limit for the accumulated dose ... currently unknown", and no publication reports the LEAF irradiation or the Hall C run; the 2023 Quantum Sensors for HEP report says "published studies in this area are lacking"	unmeasured
area	the largest device run in a GeV beam is 2×2 mm ² with 8 channels [8]; the 120 mm ² strip of §3 is 30 $\times$ that, the 496 mm ² whole gap 124 $\times$	funded R&D
latching	a hard failure, not a dead time; Argonne's ${}^{241}\text{Am}$ run — a ≈ 1 μm hot spot in a 100–200 nm wire, clean plateaus — is the reassurance [5]	addressed
rate	a few Hz per channel in normal running, 30–50 kHz at $\approx 10\sigma$ from halo [6], over three orders below the $\approx 10^8$ s ^{−1} a nanowire's fall time allows	closed

Table 6 — obstacles. None is a physics objection; all are load-bearing. Cryogenics also inverts the packaging argument: windowless operation in the beam vacuum and an edge a few hundred nanometres from the substrate rim are exactly what a near-beam sensor wants, but a device that cannot see a 300 K surface ($\approx 3 \times 10^{18}$ photons cm^{−2} s^{−1} above its 0.04 eV threshold) must sit behind a cold shield, which is material in front of the sensor; biasing below the photon threshold resolves it in principle and has never been demonstrated next to a beam.

6 · The ePIC pot geometry has moved, and has been re-measured

The aperture of §1 was first measured in the September-2024 `epic-main`. The current main branch [3] tiles the pots with 16×16 mm modules in place of 32×32 , carries per-energy 10σ offsets in `beamline_*.xml` in place of an energy-independent insertion, and has the 1 mm aluminium RF shields commented out ("we don't know if we will even need it", October 2025); the implied blind block at 18×275 is 16 mm (x) $\times$ 2.7 mm (y) against 32×7 . The old block gave 32 mm / 30.6 m = 1.046 mrad against the 1.03 mrad the transport then measured, the 1.5% agreement being the check that the measurement and the file reading confirm each other. The current block, re-measured on 2026-08-28 in `epic-main` at git 9aaa2969, gives 16 mm / 29.97 m = 0.53 mrad against 0.55 measured and 2.7 mm / 2.93 m = 0.92 mrad against 0.80 – 0.95 measured: the same check, now closed horizontally at all three configurations and vertically at 18×275 and at 10×100 , where the two sides (2.50 up, 1.95 down) straddle the geometric 2.12 mrad and the fitted -0.81 mm vertical orbit offset accounts for the split. The 5×41 vertical is the one cell the check cannot reach: the 29.6 mm insertion lies past the 6.00 mrad the vertical ladder runs to, so the 6.49 mrad the code carries there is the insertion over the $R_{34} = 4.56$ m measured on 2026-08-29 in a zero-insertion copy of the geometry — an extrapolation of a measured lever, not a measured edge — and the plane is shut because 6.49 mrad is past the 5 mrad at which the far-forward routing ends.

The correction runs in opposite directions at the two ends. At 18×275 the 16 mm block *implies* 0.53 mrad against the 0.55 mrad the first hit measures — the published triple $2.50 / 1.51 / 0.53$ is the block arithmetic throughout, and it closes on the measurement to 1–4%, well inside the 0.92 mrad Yellow Report envelope — so at the published optics the beam binds there, by eight orders of magnitude — but four times the 0.12 mrad tagging envelope, so under the tagging optics the silicon binds again until a layer follows. At 5×41 the per-energy insertion runs the other way, because the pots now retract for the larger low-energy beam: it holds the three inner silicon bands off the beam plane at $|y| \geq 29.6, 27.5$ and 18.0 mm, so a recoil at $y \approx 0$ must clear 48 mm of x rather than the 16 mm at 18×275 , and the aperture there measures 2.50 mrad — 48 mm / 19.24 m — now *outside* the 2.20 mrad envelope: the silicon and not the beam is the binding constraint at that configuration, and the ordering the September-2024 reading gave is reversed at both ends. The 5×41 figure carries one condition the other two do not. `epic-main` ships two lattices for that ring setting — the separately re-matched proton one the baseline compact file uses, and a $Z/A = 0.5$ file at exactly the ${}^6\text{Li}$ fill point — and they are not a field scale of one another: $R_{12} = 19.24$ m against 29.81 and an edge of 2.50 mrad against 1.61 — the first hit lands at 1.60 mrad on $+x$ and 1.70 on $-x$, and 48 mm / 29.81 m = 1.61 is the threshold the code applies — a factor 1.55 in the lever and 0.64 in the aperture (`farforward.POT_LEVERS_LIGHT_ION_LATTICE` , `reco.RP_APERTURE_MEASURED_LIGHT_ION_LATTICE`). A ${}^6\text{Li}$ fill at 81.6 GV is taken here to run in the baseline file at twice the field — the same magnets and the same orbit — which is an educated guess, stated as one, with the $Z/A = 0.5$ file carried as the labelled alternative; the far-forward working group is asked to confirm it, and the difference between the two is the size of the systematic on every horizontal 5×41 millimetre in this report — the alternative file carries no vertical lever, its ladder having been run through the real 29.6 mm insertion, which shuts that plane, so the vertical half of Table 5's 5×41 separation is left open rather than bounded. The re-measurement itself is done (`plans/09 B1`, `tools/fullsim`); the curves of Figure 1 are unaffected, which is why they are plotted against the half-width, and only the marker labelled silicon moves. R_{12} , R_{34} and D are measured per configuration since 2026-08-28, so millimetres are quoted at all three — in the published lattice: the millimetres this report quotes at the tagging optics carry that transport through the de-squeeze unchanged, which is the phase-advance question §7 puts to C-AD (Report 1 §6.1).

7 · Recovering the tag: beam and detector

The levers available at IP6, priced at 5×41 against the Yellow Report baseline, are unequal. A pot standoff of 8σ rather than 10σ is worth $\times 230$ at no luminosity, against a halo rate and an aperture hierarchy that do not yet exist and a RHIC precedent of 8σ for the pot with the active edge at 10σ . Detector relocation along the far-forward line is worth $\leq \times 1.75$. A lower beam energy is exhausted at $\gamma = 43.5$; a colder beam from lower intensity runs backwards, since fixed-rigidity injection puts an $A/Z = 2$ ion at half the proton's γ with $\approx 2.5\times$ the space charge and there is no store cooling in the baseline to settle into; dispersive tagging at IP6 is 30 – $60\times$ short of the 0.1 rigidity threshold of Gamage et al. [12] and the beam pipe cannot be instrumented past ≈ 30 m because of the crab cavities; and the LHC's " 3σ " is a pot window with the silicon at 4.8 – 5.4σ in fills of 3×10^{11} protons, not an active edge. The lever that closes is β^* , and Report 1 §6.1 prices it in the form

this report has used throughout: the horizontal plane alone, the pots following, at $r_h = 46, 164$ and 89 and $1/6.8, 1/12.8$ and $1/9.5$ of the luminosity, yielding $2.4, 2.4$ and 6.1×10^6 tagged recoils per year — $2.3, 2.4$ and 6.0×10^6 through the chain of Table 2 — an $8\text{--}11\sigma$ measurement of the shape term per year (quoted on the anchor's a_2 footing; the observed $\cos 2\phi$ modulation is $2a_2$, so the significance is conservative by $\times 2$ — Report 1 §6.3) and 5σ on a 1% exotic-glue term in 3.0, 5.5 and 2.6 years, both on the counting floor of the best (x, Q^2) super-bin alone, against $\delta a_e = 0.00121, 0.00111$ and 0.00074 per year on the flat term from the seven- $|t|$ -bin fit of the whole tagged sample — each projection assumes the full luminosity of the year at its own optics and isotope, so a run plan that gives this channel a share f of the year multiplies those years by $1/f$ (plans/07 WP2). A second interaction region's secondary focus at $\approx 20\%$ efficiency — our interpolation between the published d, ^3He , ^4He and ^7Li efficiencies, no ^6Li case being published — would be worth $4\times, 10\times$ and $6\times$ that yield at equal luminosity, a luminosity that is not published.

The energy ordering is now the reverse of what the earlier analysis found, and for the same reason: at the published optics the low-energy configuration led by five orders because a fixed 0.727 mrad envelope is the smallest fraction of the 41 GeV/u recoil's angular scale, while at the tagging optics each configuration's envelope shrinks with its own divergence and the top configuration, with four times the coherent events at equal luminosity, leads. Both orderings are carried by the same proton-derived placeholder — $10 \text{ fb}^{-1}/\text{u}$ per year taken as the high-acceptance luminosity at every configuration (Report 1 §3.1) — applied to a lithium beam; the Yellow Report's own per-configuration proton luminosities differ by up to $23\times$ (Report 3 Table 3), and with a species-limited lithium luminosity the ranking is open. The two questions for C-AD are those of Report 3 Table 9: can IR6 be matched with the horizontal hadron β^* raised $46\text{--}164\times$ and the electron β^* co-de-squeezed so ξ_p stays ≤ 0.015 , the Conceptual Design Report's hadron beam-beam design limit [15], with what IP-to-pot phase advances; and what physically sets the 10σ retraction. A third belongs with them, and it is a scheduling question rather than a lattice one: the de-squeeze that the ^6Li tags need costs the ^7Li tags $\times 1/6.8$ to $\times 1/12.8$ in luminosity if they share the ^6Li optics — a net factor $7.7\text{--}14.5$ at a ^7Li -only optimum (§2.1) — and buys them nothing, so the two isotopes are separate runs: the request is for two run configurations — the tagging optics for ^6Li , the published one for ^7Li — and therefore for one new optics, not two. The published high-acceptance β^* factors shrink toward low energy — 5.21 at 18×275 , 1.49 at 10×100 , 1.00 at 5×41 , where the two optics tables are identical — so the headroom the designers found runs out where the smallest absolute ask, $\beta_x^* \approx 42 \text{ m}$, sits.

8 • Conclusion

The far-forward lithium tags are an angle, and the angle is set by the machine before it is set by the sensor. At the published optics of 10×100 and 18×275 the 10σ envelope sits inside the silicon measured in the current geometry, and at 5×41 the silicon sits outside it, so no closer approach by any technology tags a coherent recoil at more than 7×10^{-8} — the 5×41 envelope, the most permissive of the three numbers a published optics allows — or more than the 1.5% off-rigidity slice of the α spectators together with the $0.1\text{--}1.0$ point over-rigid branch on the inner side of the bend, $1.9 / 1.7 / 2.6\%$ in all, and the earlier finding that a closer approach was worth $\times 26$ to $\times 569$ was an artefact of a proton divergence applied to a lithium beam. Under a lithium tagging optics the envelope shrinks by $6\text{--}9\times$ and the silicon, unmoved, becomes the whole limit: a layer that follows the envelope tags $25\text{--}37\%$ of coherent recoils and $22\text{--}31\%$ of α spectators at $1/6.8\text{--}1/12.8$ of the luminosity, gives seven populated $|t|$ bins per configuration through the full reconstruction chain, and its gain lives in a strip $50\text{--}180 \mu\text{rad}$ wide at the inner edge. The layer and the optics are multiplicative, and neither is worth anything alone.

Superconducting nanowires appear nowhere in that layer. The near-beam role fails because a sensor's dead edge is worth 0.005 mrad against a gap set by module tiling and by a rule that protects the machine, and the charge-identification role fails because the incumbent AC-LGAD already digitises an 8-bit charge across four planes, because the nanowire's one bit is nearly sufficient in principle but its 50% geometric fill cannot reach 95% efficiency, and because the $\alpha + d$ breakup is two hits tens of pixels apart. That is a conclusion about this application, not the technology: the field result that makes it unique — saturated efficiency to 5 T parallel to the device plane [10] — is worth nothing in a field-free drift and everything inside a magnet. What would help the lithium tags, in order: the tagging optics; the x -moving near-beam layer ePIC is already designing, following it; the charge EICROC already digitises, with its dynamic range published; two-hit topology for the $\alpha + d$ channel, worth an 84% veto at the tagging optics and next to nothing at the published ones, subject to how far the pot stations extend; and the IR-8 secondary focus.

The optics carries the tagged observables too, and there it buys reach rather than rate. At every published optics the ${}^6\text{Li}$ α tag admits nothing below a spectator momentum of 0.15 GeV/c — the surviving events are the off-rigidity high- k tail, median 0.32 GeV/c — so the tagged tensor asymmetry is a one-point measurement at the $O(1)$ end of the S/D interference; under the tagging optics the near-beam tail reopens and a third to a half of the accepted sample falls below 0.15 GeV/c, at little cost in tagged events per year — the acceptance gain outruns the luminosity loss at 5×41 and 18×275 , by $1.8\times$ and $1.2\times$, and falls 19% short of it only at 10×100 . ${}^7\text{Li}$ is the counterexample that fixes the shape of the whole result. Its α is off rigidity, inside the momentum window, and tagged at 96–99% in angle whatever the envelope does — at 5×41 the measured per-band insertion leaves 57%, station 2 only, on a module edge (§2.1) — so the de-squeeze buys it $\times 1.02$ in acceptance for $\times 1/7.9$ – $1/14.8$ in luminosity at a ${}^7\text{Li}$ -only optimum, $\times 1/6.8$ – $1/12.8$ if it shares the ${}^6\text{Li}$ optics (§2.1): for ${}^7\text{Li}$ the tagging optics is a net loss of a factor 7.7–14.5, and a near-beam layer is worth nothing. Every conditional in this report has the same structure — the near-beam question is only ever a question about beam-blind tags — and the two isotopes therefore want different machine optics and are different runs.

References

- [1] R. Abdul Khalek et al., *Science Requirements and Detector Concepts for the Electron-Ion Collider: EIC Yellow Report*, Nucl. Phys. A 1026 (2022) 122447, arXiv:2103.05419, §14.5.3 (refs/2103.05419_part4.pdf): the four EIC SNSPD applications; edgeless operation "within a few 100 nm of the substrate edge"; and Tables 10.1–10.2, the beam divergences per configuration and optics.
- [2] W. Armstrong (PI) et al., *Superconducting Nanowire Detectors for the EIC*, FY22 EIC-Related Generic Detector R&D proposal #18, §3.3 Concept 1, reviewed and partially funded (Report of the 2022 EIC-related Generic Detector R&D Review Committee, June 2023).
- [3] eic/epic, branch main: compact/far_forward/roman_pots_eRD24_design.xml and compact/fields/beamline_{18x275,10x100,5x41}.xml, read directly 2026-08-26. Module 1.6×1.6 cm; per-energy offsets with the comment "These are the ten-sigma cuts for the Roman pots, translated to the physical layout we currently have. They are not perfectly ten-sigma for reasons of physical geometry."; RF-shield components commented out, "Oct. 2025".
- [4] S. Lee, T. Polakovic, W. Armstrong, A. Dibos, T. Draher, N. Pastika, Z.-E. Meziani, V. Novosad, *Beam Tests of SNSPDs with 120 GeV Protons*, Nucl. Instrum. Meth. A 1069 (2024) 169956, arXiv:2312.13405 (refs/2312.13405.pdf): 12 nm NbN, 30×30 μm^2 , wire widths 300–800 nm at 2.82 K; Eqs. 1–2; $r_s = 134$ nm; the ≈ 250 nm optimum; background rising exponentially above $I_b/I_c = 0.8$.
- [5] S. Lee, W. Armstrong, J. DiPreta, C. Dulya, V. Novosad, T. Polakovic, *Optimization of Cryogenic Detector Test Station by Rejecting Electromagnetic Interference*, Nucl. Instrum. Meth. A 1093 (2027) 171953, arXiv:2601.03158 (refs/2601.03158.pdf): the 5.5 MeV ${}^{241}\text{Am}$ α at a ≈ 1 μm hot spot; clean plateaus with no latching. And A. Jentsch, ePIC Collaboration Meeting, JLab, July 2025: the x- and y-moving layers.
- [6] ePIC Collaboration, *Preliminary Design Report*, 29 September 2024 (the 1 October 2024 snapshot of chapters 2 and 8, doi:10.5281/zenodo.13866213, p. 203), the Roman Pots and off-momentum detectors section: "Rates during normal operations, with expected vacuum of 10^{-9} mbar, are a few Hz/channel. However, the beam halo could potentially provide rates of 30–50 kHz at $\sim 10\sigma$ from experience of Roman pots at STAR."
- [7] C. Wang et al. (Caltech / JPL / Fermilab), *Towards High-Efficiency Particle Detection Using Superconducting Microwire Arrays*, arXiv:2510.11725 (refs/2510.11725.pdf): 8-channel 1×1 mm^2 WSi, 1 μm wires on a 3 μm gap (25% fill), 0.8 K; 75% fill-normalised efficiency and 130 ± 17 ps at CERN SPS H6.
- [8] C. Peña, C. Wang, S. Xie et al., *Characterization of a Superconducting Microwire Single-Photon Detector Array for Charged-Particle Detection*, JINST 20 (2025) P03001, arXiv:2410.00251 (refs/2410.00251.pdf): 2×2 mm^2 , 8 channels, 30 ± 1 μm spatial resolution; amplitude distributions indistinguishable across species; edge sharper than the 10 μm telescope.
- [9] R. Cristiano, M. Ejrnaes, A. Casaburi, N. Zen, M. Ohkubo, *Superconducting nano-strip particle detectors*, Supercond. Sci. Technol. 28 (2015) 124004; Suzuki et al., Rapid Commun. Mass Spectrom. 24 (2010) 3290; US Patent 8,872,109 B2; J. J. Renema et al., arXiv:1301.3337.
- [10] T. Polakovic et al., *Superconducting nanowires as high-rate photon detectors in strong magnetic fields*, Nucl. Instrum. Meth. A 959 (2020) 163543, arXiv:1907.13059 (refs/1907.13059.pdf): saturated efficiency to 5 T parallel to the device plane, ≈ 0.5 T perpendicular; 10^7 counts/s measured.
- [11] EICROC0 characterisation and the ePIC AC-LGAD specification: A. Jentsch, DIS 2023 (refs/ePIC_far_forward_talk_DIS_2023_v2.pdf); the EICROC TDC and SAR-ADC figures as given in Report 3 §6.1, attributed there to that talk and to the ePIC Collaboration Meeting, JLab, July 2025; no OMEGA/JJCLab document is held in refs/.
- [12] B. R. Gamage et al., *Design Concept for the Second Interaction Region for Electron-Ion Collider*, Proc. IPAC'21, Campinas, paper TUPAB040, pp. 1435–1438, doi:10.18429/JACoW-IPAC2021-TUPAB040, arXiv:2105.13564, Table 1, the row "Minimum $\Delta(\text{Bp})/(\text{Bp})$ allowing for detection of $p_T = 0$ fragments": 0.1 at EIC IR 1 and 0.003–0.01 at EIC IR 2.
- [13] This repository: Report 1 §6.1 and evgen/scripts/tagging_optics.py (the tagging optics priced); Report 2 §4.5 (the coherent chain); Report 3 §4–5 (the divergence derived, the tags per optics); plans/09_nearbeam_nanowire_far_forward.md and plans/10_beam_divergence_light_ions.md.

- [14] W. Cosyn, C. Weiss, *Semi-inclusive DIS on a polarized spin-1 target*, I and II, Phys. Rev. C 114 (2026) 025201 and 025202, arXiv:2603.23699, 2603.23700 (refs/2603.23699.pdf , refs/2603.23700.pdf): Part II, "Tensor-polarized spin asymmetries of order unity are achieved for spectator momenta ≥ 300 MeV, which select configurations with large D-wave".
- [15] F. Willeke et al. (J. Beebe-Wang, ed.), *EIC Conceptual Design Report*, BNL-221006-2021-FORE (2021), doi:10.2172/1765663, read 2026-09-06: §1.2 p. 5 and §4.6.1 p. 392, the hadron beam-beam design limit $\xi_p \leq 0.015$ (electrons $\xi_e \leq 0.1$), from RHIC polarized-proton experience; for ions §1.4 p. 14 scales it with Z and A, and Appendix A §A.2.2 p. A-12 caps the two IRs together at 0.03. Report 3 [11] carries the fuller pointer set.

Appendix A · Revision history

DATE REVISION

- 2026-09-06 Sourcing pass. §7's $\xi_p \leq 0.015$ is cited to the Conceptual Design Report, [15] (plans/04 #22). Reference [12] carries the paper's actual title — *Design Concept for the Second Interaction Region*, IPAC'21 TUPAB040 — and quotes the Table 1 row the 0.1 comes from; that table supports a 10–33× band, so the 30–60× and the $\leq \times 1.75$ of §7 remain under plans/04 #24, unsourced. [6] gains its public locator.
- 2026-09-02 Consistency review of the series (docs/consistency_review_2026-09-02.md), 39 findings on this report. Abstract and §1: the two-hit separation reads 3–41 mm, the α tag 22–31% as in Reports 0 and 3, and the exponential in the closest approach is the coherent tag alone (F015, F051, F231); the tagging-optics envelope cites [13] and Cosyn and Weiss enter as [14] (F247, F248); Figure 1's two 10×100 coherent markers are noted as sitting below the panel floor (F249). Table 1: the tagging-optics α row is 0.31, the caption is dated to the 2026-08-29 re-run, and its account of the gap to Report 3 Table 6 is corrected — the near-beam band is tagged only outside the rectangle, so a wider vertical admits marginally fewer and is worth under 0.01 points, and what separates the two tables in the last digit is the Monte-Carlo draw, this scan's 2×10^5 spectators from a fresh seed against Table 6's 4×10^5 (F051, F109, F050, F104). §2: the coherent tag at the silicon edge is of order 10^{-5} rather than below it, the published-optics ceiling carries the over-rigid branch beside the 1.5% slice, and a published optics allows three numbers and not six (F239, F106, F114). Table 2 and Figure 2: a_t is quoted at each bin's t_{ref} , four of the seven bins extrapolating past the anchor, and the 1.5×10^{-5} silicon row is the divergence-smeared reading on 29 recoils, $\pm 19\%$ (F074, F251, F240, F253). §2.1: the γ^2 target-mass term is the largest in the programme rather than the only one that survives, against $\leq 0.6\%$ and $\leq 2.5\%$ at the sweet spots of those two configurations; the ${}^7\text{Li}$ α tag is flat in angle, the measured per-band insertion leaving 57% at 5×41 ; the net loss is 7.7–14.5; the low-energy verdict is Report 2's before its 2026-08-28 rewrite; and the years to 5σ are the best super-bin's counting floor, with the seven-bin δa_e beside them (F054, F009, F216, F097, F007). §3, §4 and Table 6: the tagging envelope at 18×275 is 3.52 mm and $R_{12} = 29.97$ m is assumed to survive the de-squeeze (F078, F246); $r_s = 102$ –120 nm (F250); the partner veto at the measured outer edge is 0.79 / 0.82 / 0.84 (F143); the cold-stage budget is restored, the two upper configurations at 11× and 5× a pulse tube's 4.2 K lift and 41 GeV at half of it, the whole gap 496 mm² and 124×, the rate margin over three orders (F140, F117, F243). §6: 29.6 mm is the vertical insertion, which lifts the inner bands off the beam plane; the geometry check closes horizontally at all three configurations and vertically at two; the $Z/A = 0.5$ systematic is on the horizontal millimetres (F093, F141, F245). [11] attributes the EICROC figures as Report 3 §6.1 does, no OMEGA/IJCLab document being held (F227), and the foot prices the 40–60 GeV⁻² band around $B = 50$ (F175). Appendix A: the merge scale reads 21, 12 and 8.1 mm on levers of 4.2×, 6.3× and 10.2×, the single-lever separations are marked as such, the 2026-08-28 row's δa_t pair is stamped as the record of that day with the per-configuration pair beside it, and the 2026-08-29 row now records the Table 2 and Figure 2 re-derivation it had claimed only for Table 1 and Report 3 Table 6 (F113, F116, F110, F111). Figure 2 regenerated: its title is scoped to the pots-following runs and its abscissa labelled t_{ref} (F252). Left to the author: what produces §7's $\times 1.75$ relocation lever and its 30–60× dispersive-tagging shortfall (F254), and what sources the $\xi_p \leq 0.015$ it repeats from Report 3 Table 9 (F226). Wording suggestions of the same review applied: F039, F040, F073, F076, F095, F096, F112 and F151 — the shape-term significance is labelled as the anchor's a_2 footing against Table 2's modulation coefficient $c_2 = 2a_2$; §7's yield triple is Report 1 §6.1's analytic 2.4, 2.4 and 6.1×10^6 with the chain's 2.3, 2.4 and 6.0 beside it; the year's luminosity is named as the flat 10 fb⁻¹/u placeholder against the Yellow Report's per-configuration values; the ${}^7\text{Li}$ cost is separated into the ${}^7\text{Li}$ beam's own optimum and the price of sharing the ${}^6\text{Li}$ optics; the $+0.1\sigma$ at 18×275 is marked as one draw against Report 2's twenty-experiment mean; the $|t|$ bins are populated rather than clean; the 2026-08-29 closure is bounded to its bin; and the IR-8 $\approx 20\%$ is our interpolation between the published d, ${}^3\text{He}$, ${}^4\text{He}$ and ${}^7\text{Li}$ efficiencies.
- 2026-08-29 Verification pass over the round. The abstract and the conclusion follow Table 2 — seven $|t|$ bins, not four — and Table 1, whose α tag was a 28–36% one against the 28–35% of Report 3 Table 6, a different routing. Later the same day the per-configuration transport became the baseline for the tagging optics as well, and both tables were re-run on it: Table 1 and Table 6 both read 22–31%, agreeing to a tenth of a point at every configuration: the vertical half-widths at which the two are evaluated differ, but on a common sample they are worth under 0.01 points and in the downward direction, and what separated the tables was the Monte-Carlo draw. Table 2 and Figure 2 were re-derived on the same transport: the pots-following rows read acceptance 0.36 / 0.25 / 0.32 against 0.41 / 0.32 / 0.32 and N_{tag} 2.3 / 2.4 / 6.0 against 2.5 / 2.9 / 6.0 $\times 10^6$ per year, with $\delta a_t = 0.0023$ –0.0130 in the four bins below 0.08 GeV² and 0.0148–0.0190 in the sparsest; §2.1 and §7 follow the table. The $\alpha + d$ separation of Table 5 and the target-mass default of §2.1 follow the same round: $R_{34} = 4.56$ m at 5×41 , measured through a zeroed-insertion scratch

geometry, takes the median separation from 25.8 to 17.3 mm, and `InclusiveKernel` carries its target-mass term by default from 2026-08-29, so §2.1's γ^2 statement is what the kernel now computes rather than what it drops. The coherent tag a closer approach could reach is stated as the achievable 7×10^{-8} to 6×10^{-27} with the silicon-edge column named separately, rather than the aperture-only 1×10^{-5} at 18×275 , where the envelope binds. §2.1 attributes the surviving 1.5×10^{-5} at 18×275 to the horizontal inner edge, the shallowest of the three in transverse momentum, and not to the vertical, which is an upper limit 7.6 standard deviations out. Table 5's caption uses the vertical lever where the vertical binds: the scale is 21, 12 and 8.1 mm, 4–5% above the measured minima at the two lower configurations and on them at the top, and the levers differ by 4.2 $\times$, 6.3 $\times$ and 10.2 $\times$ rather than 9 $\times$. §6 distinguishes the 0.53 mrad the 16 mm block implies from the 0.55 measured, and quotes the $Z/A = 0.5$ edge as the applied 1.61 mrad. The 2026-08-28 row no longer asserts that the machine binds at every configuration, which the re-measurement later that day withdrew. Money plot 4 re-run with no inclusive b_1 in the tagged kernel: in the impulse approximation the deuteron's b_1 is the k-integral of the spectator distribution the sampler already draws from, so the inclusive term counted the S/D interference twice — invisible with the toy shape, a +0.04 offset once the digitized b_1 replaced it. The α tag is 0.0247 / 0.2545 at 10×99.5 , the folded asymmetry at $k \approx 0.33$ GeV/c reads +0.49 / –0.07, and at 8×10^6 events it closes on the wave-function prediction in that bin to 0.005 and 0.003.

2026-08-28 The published coherent $|t|$ window widened from four bins to the seven of `recopseudo.T_EDGES_PUBLISHED`, 0.017–0.25 GeV² (Report 2 §7), and Table 2 and Figure 2 were re-derived on it: every pots-following row returns 7 of 7 bins rather than 4 of 4, with $\delta a_t = 0.0023$ –0.0074 in the four bins below 0.08 GeV² and 0.0135–0.0180 in the sparsest, on the single 18×275 lever the record then used — 0.0023–0.0130 and 0.0148–0.0190 once the per-configuration transport replaced it the next day. The acceptances, tagged yields and a_e columns are unchanged, the window being a rebinning of the same transport. The abstract and the conclusion, which had kept the four-bin count, follow the table.

2026-08-28 §7 states the run-plan share that the pricing of Report 1 §6.1 assumes: each projection takes the full luminosity of the year at its own far-forward optics and isotope, so a share f of the year multiplies the 3.0, 5.5 and 2.6 years to 5σ by $1/f$ (plans/07 WP2). No number moves.

2026-08-28 Rewritten as a paper on the per-configuration Yellow Report divergences of Report 3 §4.1 and the lithium tagging optics of Report 1 §6.1, replacing the single proton-derived 72.7 μ rad envelope that every number before 2026-08-27 used. The verdict on the near-beam layer changes: at the published optics a closer approach buys nothing — on the reading of that morning because the machine binds at every configuration, and on the re-measurement later the same day because the machine binds at 10×100 and 18×275 and the silicon at 5×41 (the earlier $\times 26$ and $\times 569$ gains are withdrawn either way), while at the tagging optics a layer that follows the envelope is the whole difference between no tag and a 25–37% tag — the two are multiplicative. Figures 1–2 and Tables 1–2 re-derived (`nearbeam_aperture_scan.py`, `nearbeam_reach_gain.py`, rectangular envelope per fragment azimuth); the sizing strip of §3 re-derived at the tagging envelope; §7 restated as the pricing of Report 1 §6.1; the energy ordering reversed with the reason given. The charge-identification material of §4, the obstacle table and the geometry correction of §6 are unchanged in substance and condensed. §4.2 gained the two-hit study it had only sketched: both fragments of the breakup are sampled from one relative momentum (`spectator.breakup_lab_kinematics`, `nearbeam_two_hit.py`), which gives the topology fractions and the partner-fragment veto of Table 5 and supersedes the two lower separation rows, which had carried the retired energies; the separation itself is the angular lever plus the pot dispersion, which does not cancel between the fragments — their rigidities move in opposite directions with k_z — so the three medians are 15.1 / 18.5 / 38.4 mm and a merge is rare rather than impossible. The Yellow Report veto column is measured on 1.2×10^7 breakups per configuration, its 10×100 entry having rested on four events. $\beta = 0.30$ GeV is defined where it is first used. §2.1 gained the tagged observables: `money_tagged_azz.py` and `tagged_polarimetry_7li.py`, whose figures were still routed through the retired proton-derived 73/164 μ rad pair, were moved onto the per-configuration optics with the spectator's own lab azimuth against the rectangular envelope (plans/09 B2, B3), and the α tag they quote is the tag itself rather than the accepted fraction inside the scripts' $k < 0.6$ GeV/c histogram, which read 9–11% low. The result is a reach statement — at the published optics the ${}^6\text{Li}$ α tag admits nothing below $k = 0.15$ GeV/c, at the tagging optics half the accepted sample is there, for a 7% cost in tagged events per year at 10×100 and a 1.2–1.9 $\times$ gain at the other two — the figures of the single-lever pricing then in force, which the per-configuration transport of the row above re-priced to 36–48% below $k = 0.15$, a 19% cost at 10×100 and gains of $\times 1.8$ and $\times 1.2$ (§2.1, §8) — and a ${}^7\text{Li}$ net-loss statement, now also in §7 and the conclusion. The money-plot-4 overlay was corrected with them: it had compared an acceptance-sculpted sample against an analytic curve drawn at $\theta_k = 90^\circ$, and the ± 0.5 swing it showed between its two optics at $k \approx 0.3$ GeV/c was that sculpting rather than the wave function. A review of the routing itself found the 10σ envelope falling back on its inscribed circle wherever $\sigma_h = \sigma_v$, which discarded the azimuth at the 180/180 and 92/92 μ rad Yellow Report optics: Table 5's Yellow Report fake-rate and veto columns at 10×100 and 18×275 fall to 0.0001 / 0.02 and 0.0007 / 0.25, and Table 6 of Report 3 with them. Late the same day the pot aperture and the pot optics were re-measured in the current ePIC geometry (`epic-main` at git 9aaa2969, plans/09 B1): the silicon edge is 2.50 / 1.51 / 0.53 mrad against 2.00 / 1.35 / 1.03, so the coherent tag at the silicon runs 9.4×10^{-10} / 2.0×10^{-19} / 1.2×10^{-5} and the verdict of §2 is no longer one statement — the silicon binds at 5×41 and the machine at the other two, where the earlier reading said the opposite. (R_{12} , R_{34} , D) are measured per configuration for the first time, which cuts the $\alpha + d$ separations of Table 5 by 28 to 42%, from 15.1 / 18.5 / 38.4 to 10.9 / 10.7 / 25.8 mm, and multiplies the merge column by 10–25, and the pot's outer edge is measured at 2.85 / 3.85 / 4.00 mrad rather than assumed at 5, which costs the partner veto five points at 5×41 and

nothing at the other two. The single |t| bin the silicon row of Table 2 used to return at 18×275 is retired by a minimum-count guard, `MIN_TAGGED_PER_BIN` = 1000 against the 231 recoils a year it holds.

2026-08-27	§7 replaced by a recovery study at the γ-matched energies: the lever table, the three levers dropped with reasons (the cold-beam argument runs backwards; dispersive tagging at IP6 is 30–60× short and the pipe cannot be instrumented past ≈ 30 m; the LHC 3σ is a pot window, the right precedent RHIC's $\beta^* = 22$ m stores at 8σ). Earlier the same day: the charge argument restated as a sampled-Landau fake rate at matched efficiency (one bit costs 1.4×, a truncated mean 86× to one bit; the nanowire loses on fill factor), the $\alpha + d$ separation re-quoted from the repository's own density (median 10.9 / 30.1 / 73.4 mm), single-fragment loss identified as the dominant one-hit topology, EICROC's dynamic range promoted to the top of the unknowns.
2026-08-26	First version, revised the same day against an adversarial review that changed the sign of the charge-identification verdict: the incumbent AC-LGAD already digitises an 8-bit charge over $30\text{ }\mu\text{m}$ across four planes; $r_s = 134\text{ nm}$ is a model parameter, the counting volume differs by 5000 between Argonne's papers, the ^{241}Am α tests $\sqrt{Q}$ across β rather than z^2 at fixed velocity; the dead-edge argument retired at 0.005 mrad; the granularity argument redirected to ePIC's x-moving layer; the $\alpha + d$ topology added. Original content: the aperture scan and the chain at two apertures, the firing-threshold model, the obstacle table, the geometry correction, the question list.

All coherent acceptances use $B = 50\text{ GeV}^{-2}$, the centre of the 40–60 GeV^{-2} band of Report 1 §6.2: across it the fixed-aperture coherent entries of Table 1, exponential in $B|t|$, move by factors of 55 at 18×275 to 10^{10} at 10×100 , and the factors §2 quotes with them, while the tagging-optics row moves by $\pm 12\%$ and the ordering of the three apertures not at all; the near-beam envelope is a rectangle $10\sigma_h \times 10\sigma_v$ in angle at the Yellow Report divergences of Report 3 §4.1 (with the species step at 18×275) or at the tagging optics of Report 1 §6.1, applied to each fragment's azimuth; the measured silicon aperture is from `tools/fullsim` in the current `epic-main`, git 9aaa2969, measured 2026-08-28 (§6). Energy deposits are Bethe means in a 12 nm NbN film; the Landau most-probable value is not quoted because $\xi/l \approx 0.002$ puts the film outside the $\xi \gg l$ regime. Hot-spot radii are anchored on Argonne's 134 nm and scaled as z . Built by `reports/build_report.py`; figures from the `evgen/scripts` named in the text; display mathematics typeset at build time.