

Simulation and reconstruction chain of the inclusive and coherent $\cos 2\phi$ measurements with a tensor-polarized ^6Li beam at the EIC: assumptions, methodology and closure

C. Peng and J. Zhou

Physics Division, Argonne National Laboratory

Writing assisted by Claude (Anthropic)

2026-08-24 · rewritten 2026-08-28 after a full review of the toolchain, restated 2026-08-29 on the per-configuration transport, revised 2026-09-02 on a cross-document consistency review, re-sourced 2026-09-06 (Appendix A) · companion to Report 1 (the projections) [11] and Report 3 (the machine and the detector) · computed with `polligen.reco`, `polligen.hfs` and `polligen.recopseudo` on the PYTHIA 8 hadronic final state of `tools/pythia8`; every number is reproducible with the commands of `docs/reproduction_manual.md`

Abstract. The double-helicity-flip structure function $\Delta(x, Q^2)$ of ^6Li is projected in Report 1 as the $\cos 2\phi$ modulation of inclusive DIS on a transversely tensor-polarized beam, and as a modulation of the recoil azimuth in coherent diffraction with the ^6Li detected intact. This report describes the simulation and reconstruction chain behind those projections — the polarized event kernel, the PYTHIA 8 hadronic final state, the electron-side, hadron-side and far-forward detector responses, the kinematic reconstruction, the spin-state-sorted estimator and the conversion to Δ — lists every assumption it rests on with its provenance and leverage, and reports its closure at the reconstructed level. The analysis path consumes only quantities an experiment records — the spin state of each bunch crossing, the tensor polarization and luminosity of each state, the reconstructed electron, the hadronic $E - p_z$ sum and the Roman-Pot angle pair — with simulation-derived corrections applied as functions of reconstructed variables; a 2026-08-28 review of the toolchain found one place where that was not so (the hadronic-scale calibration, now keyed on the reconstructed point), a truncation of the PYTHIA library below $x = 16/s$ (now regenerated), and two statistical defects of the estimator (now fixed), none of which changes a headline number by more than its statistical error. At $e(10 \text{ GeV}) \times ^6\text{Li}(99.5 \text{ GeV/u})$ the four sweet-spot bins reconstruct with purities 0.56–0.75 and amplitude errors $0.9\text{--}3.0 \times 10^{-4}$ in one year, the fit closing on the response-folded truth; the coherent channel closes at the lithium tagging optics of Report 1 with $2.3\text{--}6.0 \times 10^6$ tagged recoils per year, biased by at most 1.5% of the injected value — a third of the one-year statistical error — in every $|t|$ bin with more than 10^5 recoils.

1 · Scope

Report 1 presents the measurement; Report 3 the accelerator and detector it is projected on. This report is the layer between them: what the simulation generates, what it pretends the detector records, what the analysis is allowed to read, and how well the chain returns what was put in. Section 2 walks the chain stage by stage with the boundary between generator truth and measured quantities made explicit; Section 3 collects the assumptions; Section 4 gives the methodology of each step; Section 5 reports the reconstructed-level closure; Section 6 the systematic effects that were measured on the chain and the specifications they imply; Section 7 what the chain does not model. The audit of the code that produced this revision is summarized in Appendix A and recorded in `docs/code_review_2026-08-28.md`.

2 · The chain

The observable is a Fourier coefficient of a *ratio* of spin-sorted azimuthal distributions. For an unpolarized electron on a spin-1 target with alignment axis at polar angle θ_S and azimuth ϕ_S to the beam — coincident with the virtual-photon axis in the massless limit the formula is written in, the $O(\gamma^2)$ difference measured in §4.1 — the Hoodbhoy–Jaffe–Manohar cross section of the projection state m is [1,2]

$$\frac{d\sigma^{(m)}}{dx dy d\phi} = \frac{2y\alpha^2}{Q^2} \left[F_1 + \frac{2}{3}a_m b_1 + \frac{1-y}{xy^2} \left(F_2 + \frac{2}{3}a_m b_2 \right) - \frac{1-y}{y^2} c_m \sin^2 \theta_S \Delta(x, Q^2) \cos 2\phi \right], \quad c_m = 3m^2 - 2,$$

so that a fill with populations p_m along a transverse axis has the azimuthal distribution $1 + a_2 \cos 2\varphi'$, $\varphi' = \varphi - \varphi_S$, with

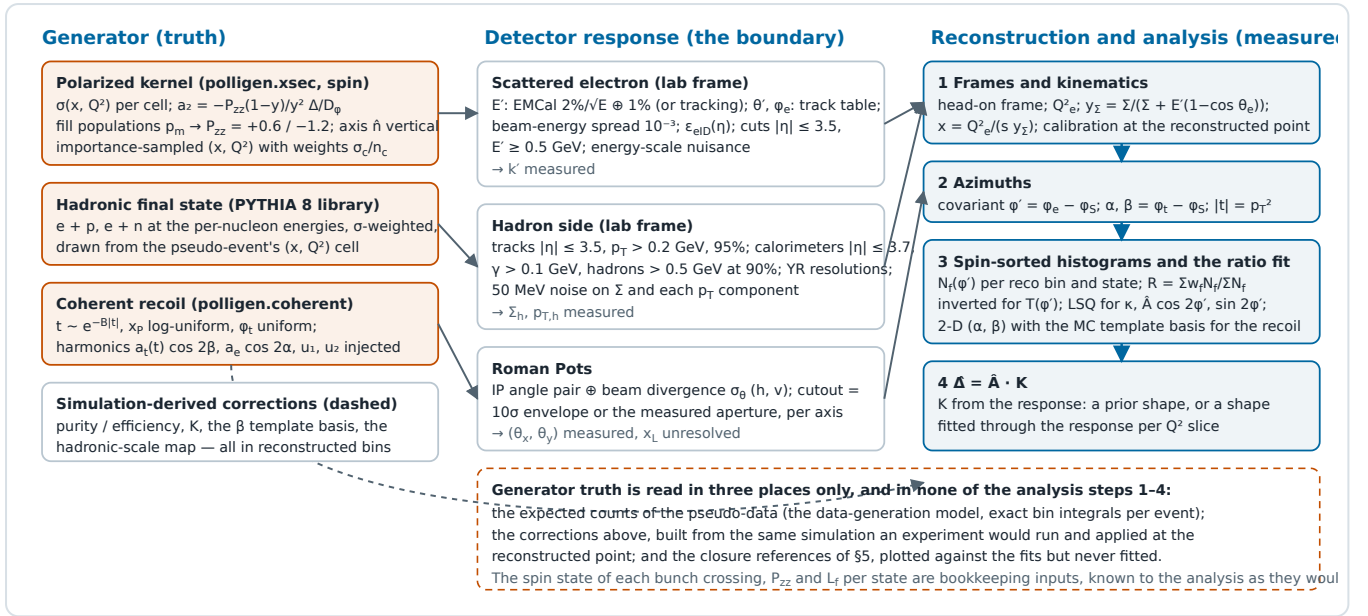
$$a_2 = -P_{zz} \frac{1-y}{y^2} \frac{\Delta}{D_\phi}, \quad D_\phi = F_1 + \frac{1-y}{xy^2} F_2, \quad \hat{A} \equiv \frac{a_2}{P_{zz}}, \quad P_{zz} = \sum_m p_m c_m.$$

The b_1 , b_2 terms are written above in the sign of this programme's own transcription of [2], the form the companion reports [11] display as well; HERMES and Cosyn et al. [12], applying the convention of [2], use the opposite one, $A_{zz} = -(2/3) b_1/F_1$ for the axis along the photon with $P_{zz} = n_+ + n_- - 2n_0$, and since 2026-08-29 so does this programme's simulation — a single constant, `polli_fastsim.asymmetries.TENSOR_LL_SIGN` = -1, adopted because it is the convention every published b_1 is quoted in. It affects only the sign of A_{zz} relative to b_1 , and with it the sign of the by-product κ of §4.3 and of the $O(\gamma^2)$ leakage of §4.1; no Δ depends on it, and it reaches a magnitude only through the 10^{-3} of the rate the b_1 term carries — one printed digit, Report 1's third-spot one-year δA in Table 4, $2.8 \rightarrow 2.7 \times 10^{-4}$.

The chain that turns this into a measurement has six stages (Schematic 1, Table 1). It begins in the polarized kernel: `polligen.xsec` evaluates the master formula, in the sign convention just stated, per (x, Q^2) cell for each projection state m of the fill, with the spin-density matrix of `polligen.spin` providing the m -populations of the two fill types and the covariant azimuth of §4.1 providing φ' ; the unpolarized cross section, F_2 and $R = \sigma_L/\sigma_T$ come from `polli_fastsim.structure` and the Δ model from `polli_fastsim.delta_models`. The tensor polarization enters the inclusive cross section only as a weight on it — the $\cos 2\varphi'$ amplitude above together with the φ -independent rate shift of the b_1 , b_2 sector, of order 10^{-3} of the rate and fitted as κ in §4.3 — and not the hadronic final state, so the hadronic system of each pseudo-event is taken from an unpolarized PYTHIA 8 [20] library of $e + p$ and $e + n$ events at the same beam energies (§4.2), weighted by their cross sections.

Three detector responses follow, all applied in the detector frame. The scattered electron is smeared with the calorimeter, tracking and electron-identification models of §3; the hadronic final state passes through a hadron-side response — tracker and calorimeter coverage, thresholds, efficiencies, Yellow-Report resolutions, calorimeter noise; and the coherent recoil is measured by an emulated Roman-Pot system that returns the angle pair at the interaction point smeared by the beam divergence and cut by the near-beam envelope. Reconstruction then restores the head-on frame and takes Q^2 from the electron, y from the Σ method on the hadronic $E - p_z$ sum, x from the two (the mixed method) and the azimuths from the covariant formula on the reconstructed four-vectors.

The last two stages are the measurement itself. Events are histogrammed in φ' per spin state and per reconstructed (x, Q^2) bin, and the $\cos 2\varphi'$ coefficient is fitted to the spin-state ratio, which cancels any azimuthal acceptance common to the two states; the amplitude is then converted to Δ with a bin-centering factor derived from the response, either from an assumed Δ shape or from a shape fitted through the response itself.



Schematic 1 — the chain, and the boundary between truth and measurement. Left: what is generated. Middle: the three detector responses, which consume the truth and emit measured quantities. Right: the analysis, which reads measured quantities, the bookkeeping of the run plan, and corrections derived from the simulation in reconstructed bins.

STAGE	INCLUSIVE $E' e^+ p \rightarrow E' X$	COHERENT $E' e^+ p \rightarrow E' X$ (G.S.)	TRUTH READ	CODE
Generation	importance-sampled (x, Q^2) over a 60×45 log grid with the generator window loosened ($y \geq 0.004$, $Q^2 \geq 0.7$ GeV ² , $ \eta \leq 3.8$, $E' \geq 0.3$ GeV) so events can migrate into the analysis window; 400 pseudo-events per cell, weight σ_c/n_c	$ t $ exponential with $B = 50$ GeV ⁻² , $x_p \in [10^{-3}, 10^{-2}]$, ϕ_t uniform; 6×10^6 recoils; injected harmonics of \$4.5	—	recopseudo.RecoResponse, CoherentResponse
Hadronic final state	PYTHIA 8.311 $e + p / e + n$ at the per-nucleon energies, $Q^2 \geq 0.7$ GeV ² , $\hat{m} \geq 0.5$ GeV, lepton radiation off; 8 M events over the three configurations; a library event from the pseudo-event's (x, Q^2) cell supplies the captured fraction of Σ_h and the p_T ratio	not generated (no electron side; α integrated analytically)	the pseudo-event's (x, Q^2, y) , to draw and scale the library event	hfs.HFSLibrary, HFSResponse; tools/pythia8
Detector response	e' : lab-frame smearing of E' , θ' , ϕ_e , beam-energy spread, $\epsilon_{elD}(\eta)$; hadrons: coverage, thresholds, efficiencies, resolutions, noise, in the lab frame	Roman Pots: angle pair @ divergence, rectangular cutout (envelope or aperture per axis)	the true four-vectors, to smear them	reco.smear_electron, hfs.HadronResponse, reco.rp_measure
Reconstruction	head-on frame; Q^2_e ; y_Σ (JB, DA available); $x = Q^2_e/(s y_\Sigma)$; hadronic-scale calibration at $(x_{\text{mixed}}, Q^2_e)$; analysis cuts $Q^2 \geq 1$, $0.01 \leq y \leq 0.95$, $W^2 \geq 10$, $ \eta \leq 3.5$, $E' \geq 0.5$; covariant ϕ'	$ t = p_T^2$, $\beta = \phi_t - \phi_5$ from the smeared angle pair; $\alpha = \phi_e - \phi_5$	none	reco.electron_method, hfs.hadronic_kinematics, reco.azimuth_wrt_lepton_plane

Pseudo-data	exact expected counts per spin state and ϕ' (or (α, β)) bin, from the per-event amplitude at the true kinematics and the reconstructed bin assignment, with the ϕ' -efficiency harmonic (common or per fill) and the relative-luminosity offset switched on; Poisson-fluctuated at 10 and 100 fb ⁻¹ /u		the per-event a_z and dilution (data generation)	RecoResponse.expected_counts, CoherentResponse.expected_counts_2d
Estimator	spin-state ratio per bin, inverted for T(ϕ'), weighted LSQ for κ , $\hat{A}$, $\hat{A}_{\sin}$	the same in (α, β) with the acceptance-weighted β template basis from the response, the u_1, u_2 nuisances subtracted, and three parity-forbidden sine columns as null tests	the basis and the truth references only	reco.harmonic_ratio_fit, harmonic_ratio_fit_2d
Extraction	$\Delta = \hat{A} K$, $K = \Delta_{\text{prior}}(x_c, Q^2_c)/\text{fold}(\Delta_{\text{prior}})$ or from the shape fitted through the response	$a_t(t_{\text{ref}})$, a_e per t bin	the response (fold, purity, efficiency)	recopseudo.fold, fold_shape_fit, delta_from_amplitude

Table 1 — the chain stage by stage, and where generator truth is read. The pseudo-data are generated from the truth, as data are generated by nature; the reconstruction and the estimator read measured quantities; the corrections are what an experiment takes from its own simulation.

3 • Assumptions

Table 2 lists every input the chain consumes, with its value, its provenance and what a wrong value would move. Three classes are distinguished: inputs the accelerator or the detector will fix (to be asked for precisely; Report 3 carries their full specification), physics-model inputs (scenarios, bracketed by bands in Report 1), and idealisations of this analysis (to be replaced or quantified as the chain matures). Idealisations that are structural — no radiative corrections, no backgrounds, no coherent electron side — are in §7.

INPUT	VALUE USED	PROVENANCE AND STATUS	LEVERAGE IF WRONG
Beam and spin			
⁶ Li energies	40.8 / 99.5 / 137.5 GeV/u	derived: γ -matched to the 41 / 100 / 275 GeV proton, rigidity-capped at the top [7,11]; the chain publishes at the mid configuration	every far-forward cut ($\propto A p_u$); the sweet-spot y values
beam divergence σ_θ (h / v)	220/380, 180/180, 92/92 μrad	YR Table 10.1 proton values (its strong-hadron-cooling table: the start-of-store emittance under the injection-cooling baseline, Report 3 §3.2), gamma-matched species step [7] (reco.sigma_theta_for); no lithium optics is published	coherent tag acceptance, exponentially; $\Delta p/p$ 7–10 $\times 10^{-4}$ enters only the pot dispersion term
tagging optics (coherent channel)	$\beta^*_x \times 46 / 164 / 89$, pots following the 10 σ envelope 0.36×3.8 , 0.19×1.8 , 0.12×0.92 mrad	this programme's proposal (Report 1 §6.1): the optimum of acceptance $\times$ luminosity with the vertical plane at high acceptance, $L/L_{\text{HA}} = 1/6.8$, $1/12.8$, $1/9.5$; assumes the IP-to-pot transport of each configuration's own measured baseline levers ($R_{12} = 19.24 / 21.25 / 29.97$ m, $D = 0.311 / 0.287 / 0.292$ m, plans/09 B1; farforward.tagging_optics_point , levers="per-config" since 2026-08-29) surviving the de-squeeze parallel-to-point, and pots that move	without it no recoil is tagged at any configuration (§4.5, Figure 3c)

P_{zz} of the two spin states	+0.6 / -1.2, equal luminosity	in-ring placeholder (ANL source target ≥ 0.8); the $m = 0$ -rich state at $-2P_{zz}$ assumes the same source purity; taken as exactly known by the analysis	statistical error $\propto 1/P_{zz}$; a scale error is a 1:1 normalisation on Δ (§6)
spin axis	vertical, $\varphi_S = \pi/2$, $\theta_S = 90^\circ$	the stable spin direction of the arcs, assumed to be the axis delivered at IP6 — no rotator or snake configuration is defined for a lithium fill (Report 3 §7); an azimuthal error is self-calibrated by the $\sin 2\phi'$ term ($\tan 2\delta$ recovered exactly in a test), a polar tilt dilutes by $\sin^2\theta_S$	0.8% dilution at 5° tilt; none from the azimuth
spin-state pattern	bunch-by-bunch alternation	required, not assumed: the ratio cancels only an acceptance common to both states (§4.3)	a 10^{-3} difference of the $\cos 2\phi'$ acceptance harmonic between the states fakes 5.6×10^{-4}
luminosity	10 / 100 fb $^{-1}$ per nucleon	YR convention for one / ten years [7]. The 0.5/0.5 share, with a 10^{-3} offset unknown to the analysis, is between the two SPIN STATES of this measurement and divides whatever luminosity it is given; it is not the programme share of plans/07 WP2, which divides the year between observables, optics and isotopes. Each projection assumes the full luminosity of the year, and both closures below are run at the full 10 fb $^{-1}$ /u	statistics; the offset enters the ratio as a constant

Central detector

EMCal σ_E/E for e'	2%/√E $\oplus$ 1%	YR requirement of the backward crystals [7], applied at every η (<code>RecoModel.emcal_eta_table = False</code>); exact at the three of the mid configuration's four sweet spots with $\eta < -2$ and optimistic by a factor 2 at the fourth, $\eta = -1.65$, where the YR asks 7%/√E (Report 3 Table 8) — those four spanning $\eta = -2.9$ to -1.6 , the twelve of the three configurations -3.4 to -0.9 , the twelfth of them in the barrel; switching the η -dependent table on changes Δ by 0.000% at the first three spots and 0.10% at the fourth, and by up to 0.7% at 5×40.8	$\delta Q^2/Q^2$ (subdominant to the angles)
tracking σ_p/p	$\sqrt{(a/p)^2 + b^2}$, η -piecewise	YR requirement (Table 11.2, and Fig. 8.3 / Table 8.20 which differ backward), endcap constants padded; a requirement, not an achievement	only with <code>energy="best"</code>
track angular resolution	5 / 3 / 2 / 1 mrad by η (3 mrad at spots 1–2, 2 mrad at spots 3–4)	placeholder — no YR angular requirement exists; it is also the only stand-in for the 81–211 μ rad beam-divergence floor on θ'	$\delta Q^2/Q^2 = 5.2\%$ at spots 1–2; a factor $\lesssim 3$ bracket (purity 0.66 $\rightarrow$ 0.70 at a third of it)
electron identification	$\varepsilon_{eID}(\eta) = 0.70\text{--}0.95$	constructed from the ATHENA [16] and ECCE proposal values but anchored at one point — ECCE's 0.95 at $\eta = -2.0$ — the barrel matching neither and the forward tail running below both; no ePIC curve [17]; no purity or background model	none on Δ beyond statistics — a fill-independent efficiency cancels in the ratio (an η tilt of 0.05 moves Δ by 0.02%)
electron energy scale	nuisance $\pm 1\%$	unknown to the analysis; $d \ln x/d \ln E' = 2 - y$ for the Σ method	0.3–1.0% on Δ per 1%
hadron-side coverage and thresholds	tracks $ \eta \leq 3.5$, $p_T > 0.2$ GeV, 95%; calorimeters $ \eta \leq 3.7$; $\gamma > 0.1$ GeV; hadrons > 0.5 GeV, 90%	stand-ins of the right magnitude in the detector frame (the ePIC nominal forward reach is 4.0; B0 at $\eta = 4.6\text{--}5.9$ not modelled)	the captured fraction of Σ_h (0.70–0.85 at the sweet spots, Table 3); 4.0 recovers a third of the forward loss (0.170 $\rightarrow$ 0.115 at the first spot) with $\delta y/y$ unchanged

hadron-side resolutions	YR tables per region	requirement values [7], with the YR 2% EMCal constant	subdominant to the noise below $y = 0.02$
calorimeter noise on Σ_h	50 MeV Gaussian per event on Σ and each p_T component	a one-parameter stand-in for the noise and threshold floor whose leverage at low y Arratia et al. show for H1/ATHENA [18]; no ePIC number exists (plans/04 #21)	the $y \approx 0.01$ resolution: $0.22 \rightarrow 0.54 \rightarrow 1.01$ at $y = 0.005$ for $0 \rightarrow 50 \rightarrow 100$ MeV
hadronic energy scale	calibrated from the simulation at $(x_{\text{mixed}}, Q^2_e)$; residual $\pm 2\%$ as a nuisance	the analysis's own calibration; keyed on reconstructed variables (2026-08-28)	$0.1\text{--}0.6\%$ on $\hat{A}$ per 1% residual

Far forward

Roman-Pot measurement	angle pair $\oplus \sigma_\theta$ (h, v); no detector term	the rigidity-window routing of the standard far-forward acceptance model [8]; beam effects dominate the ePIC far-forward resolution [15]; x_L unresolved; the recoil's own dispersion at the pots ($D x_P/R_{12} = 0.16 / 0.14 / 0.10$ mrad at $x_P = 0.01$, inside the 10σ tagging envelope of $0.36 / 0.19 / 0.12$ mrad but above its $36 / 19 / 12 \mu\text{rad}$ horizontal σ_θ) is stated rather than modelled; what the horizontal term of the tagging optics folds in is the beam's own momentum spread, which enters the same way at $17 / 13 / 7 \mu\text{rad}$	$ t $ smearing $2p_T\delta p_T$; the ϕ_t dilution the template basis absorbs
pot aperture	2.50/shut, 1.51/2.12, 0.53/0.92 mrad (h/v)	measured with an intact ${}^6\text{Li}$ in the current ePIC geometry (<code>epic-main</code> at git 9aaa2969, 2026-08-28; <code>tools/fullsim</code>), on 0.05 mrad angular ladders at four azimuths; at 5×41 the vertical plane admits nothing and the code carries 6.49 mrad, the 29.6 mm insertion over the $R_{34} = 4.56$ m measured on 2026-08-29 in a zero-insertion scratch geometry ($R_{12} = 19.18$ m against the real 19.24 as its control), and the horizontal 2.50 mrad stands on the ePIC baseline proton lattice for that ring setting — the $Z/A = 0.5$ file at the ${}^6\text{Li}$ fill point gives 1.61 mrad (48 mm / 29.81 m, against a first hit at 1.60 on $+x$ and 1.70 on $-x$), a $\times 0.64$ systematic on that configuration alone	with the pots fixed at this edge the tag is 10^{-5} – 10^{-19} , and 9×10^{-10} – 6×10^{-27} at the published optics, where the 10σ envelope binds at the two upper configurations (§4.5); the tagging optics assumes pots that follow, as would the secondary focus of the second interaction region [10]
exclusivity	pure coherent sample	no breakup background, vetoes or Z-identification modelled (plans/06); the incoherent yield is m-state blind and dilutes rather than fakes	the $ t $ -fit purity of Report 1 §6.4, and the $\alpha + d$ partner-fragment veto, which reaches 0.82–0.84 of the α fakes at the tagging optics, 0.79–0.84 at the measured outer edge of the pots (Report 4 §4.2)

Physics model

F_2, F_1, R	toy F_2 ($\pm 37\%$ of CT18 over the acceptance); $R = 0.18/(1 + Q^2/50)$	placeholders with PDF-grid and R1998 backends behind the same hooks (Report 1 §3.2); F_2 cancels exactly in the amplitude, R does not	rates and statistical errors to $\times 1.5$; Δ/F_1 by $+18 / +18 / +8 / -4\%$ at the four spots if the R1998 world fit [19] replaces the placeholder
$\Delta(x, Q^2)$	moment-constrained ansatz A , dilution 1/3; α_s from the CT18NLO table	Report 1 §4; whether Δ should carry the rank-2 α -d transfer 0.92 as b_1 does ($1/3 \rightarrow 0.307$, an 8% reduction) is left open in the Δ -model registry, and no projection here assumes it; the LO α_s fallback (used without the <code>parton</code> package, now announced) raises the normalisation by 14%	the injected amplitudes, not the closure

coherent scenario	$f_0 = 0.04$, $B = 50 \text{ GeV}^{-2}$, $\epsilon_{B0} = -0.08$, $a_e = 0.01$, $u_1 = 0.05$, $u_2 = 0.02$	Report 1 §6; $u_1 = 0.05$ is ZEUS's central A_{LT} at $0.01 < x_p < 0.1$ and $u_2 = 0.02$ lies within one ZEUS error (0.024) of zero, 1.2–1.4 σ above either central A_{TT} ; both are of the size the diffractive peak allows, where the asymmetries are consistent with zero within ± 0.07 [13,14]; taken as known by the fit	rates and the injected harmonics; δu_2 leaks into a_e at first order (§6)
$a_t(t)$ template	linear in $ t $	the anchor's low- $ t $ behaviour [9]; one coefficient per $ t $ bin	a different shape is a model systematic to be varied
radiative corrections	none in the published response; collinear ISR available as a hook (<code>polligen/radiative.py</code> , <code>RecoResponse(isr=...)</code> , default off)	bounded, not calculated: the exponentiated leading-log spectrum [21,22] as a kinematic migration. Collinear ISR fakes no ϕ' modulation and cancels in the spin-state ratio; $x = Q_e^2/(s y_z)$ is exact and only the Q_e^2 label migrates, by $1/(1-z)$. Tensor-sector RC uncalculated (plans/05 §5.5), and no unpolarized study bounds it	Δ by +0.5 to +1.2% at the four sweet spots on the Gaussian y stand-in, +0.4 to +0.9% through the calibrated PYTHIA final state (mean over eight response seeds; $\leq 2.9\%$ once the generator window is opened below $Q^2 = 0.15 \text{ GeV}^2$ and the low- Q^2 feed-in saturates); $\leq 0.25\%$ with a HERA-style $E - p_z$ window, which keeps 87% (stand-in) / 98% (PYTHIA, calibrated Σ) of the non-radiative rate and is a contingency rather than a default. §7

Table 2 — the assumptions. "Requirement" means the Yellow Report value, which is a design target and not a measured performance. Every number in this report is conditional on this table.

4 • Methodology

4.1 • Frames and azimuths

The generator works in the head-on frame (ion along +z, electron along -z); the detector sees the lab frame, in which the electron beam defines the ePIC axis and the ion enters at -25 mrad in x. The two are related by a rotation of 12.5 mrad about y and a boost $\beta_x = \sin 12.5 \text{ mrad}$; energies change by 8×10^{-5} and the vertical spin direction not at all. Every detector effect — the electron smearing, the hadron acceptance and thresholds, the Roman-Pot angles — is applied in the lab frame, and the measured four-vectors are transformed back before any kinematic quantity is formed, as an ePIC analysis does. For the electron the transformation changes only odd harmonics of the azimuth ($\approx 10^{-3}$ in $\cos \phi$, $\cos 2\phi$ untouched to 10^{-17}), so it cannot fake the observable; for the hadrons it moves the forward calorimeter edge by up to ± 0.25 units of η around the ring, which is worth +0.7 to +1.5% of the captured Σ_h at the sweet spots relative to a head-on-frame acceptance.

The azimuth of the observable is the covariant one used for every transverse-spin azimuth in SIDIS and exclusive analyses [3,4], applied to the spin four-vector S and, for the recoil, to P':

$$\cos \phi_S = -\frac{l_\perp \cdot S_\perp}{|l_\perp| |S_\perp|}, \quad \sin \phi_S = -\frac{\epsilon_{\mu\nu\rho\sigma} l^\mu S^\nu P^\rho q^\sigma}{(P \cdot q) \sqrt{1 + \gamma^2} |l_\perp| |S_\perp|}, \quad v_\perp^\mu = g_\perp^{\mu\nu} v_\nu, \quad \gamma^2 = \frac{M^2 Q^2}{(P \cdot q)^2}.$$

With the axis transverse at lab azimuth ϕ_S this reduces exactly to $\phi' = \phi_e - \phi_S$ for a massless target and to $O(\gamma^2)$ for the nucleus — at most 1.4 and 1.2 mrad at the sweet spots of the mid and top configurations ($\gamma^2 \leq 0.0057$ and 0.0049) and 5.9 mrad at those of the low configuration ($\gamma^2 \leq 0.025$), measured against the lab-angle shortcut at the twelve spots by `evgen/scripts/target_mass_bound.py` — and it is verified against an explicit collinear-frame construction to 10^{-13} . $\cos 2\phi'$ is even under $\phi' \rightarrow -\phi'$ and $\phi' \rightarrow \phi' + \pi$ and the axis is headless, so no orientation convention enters; a

misalignment δ of the assumed axis turns the modulation into $\cos 2\delta \cos 2\phi' + \sin 2\delta \sin 2\phi'$, and the fitted $\sin 2\phi'$ term returns $\tan 2\delta$ — the self-calibration of the axis. The tensor structure functions b_1 – b_4 leak into $\cos 2\phi'$ at $O(\gamma^2)$, because an axis transverse to the beam is not transverse to the photon: the exact kernel of Cosyn et al. [12] Eqs. (9), (10), (14), (17) and (24) — in the generator since 2026-08-29, off by default so that this chain's tensor sector runs on the massless path it was built on, whereas the longitudinal kernel has carried its exact finite- γ $A_{\parallel} = D_{\gamma}(A_1 + \eta A_2)$ by default since the same day, `InclusiveKernel(target_mass=True)` being a separate switch that reaches no Δ — makes the leakage an effective Δ of (0.14–0.16) $\gamma^2 b_1$, measured at the twelve sweet spots by `evgen/scripts/tensor_gamma_leakage.py` — small because its leading-twist and twist-3 channels cancel to three parts in three, leaving the twist-4 Eq. (17e) term almost alone — a cancellation that holds at the $b_3 = b_4 = 0$ the kernel defaults to, those two being unmeasured: $b_3 = b_4 = 0.1 b_2$ takes the coefficient to (0.23–0.26). At the configurations of the closures below it is -0.033% of the amplitude at worst, -0.052% at that tenth of b_2 (-0.109% and -0.175% at the low configuration of the coherent closure, where γ^2 is four times larger), it carries the sign that cancels rather than fakes, and it is subtractable in situ rather than against the A_{zz} of the same data: 99.96% of it is the b_2 term — and $b_2 = 2xb_1$ at leading twist — entering in the same combination as the ϕ' -independent constant κ of the ratio fit, the by-product of §4.3 whose sign the tensor convention adopted the same day fixes, so that $L = A_{\text{leak}}/\kappa$ is a pure kinematic function of (x, Q^2, y) and the correction reads $\hat{A} - L D_{\phi} \hat{\kappa}$ in measured quantities alone. Since 2026-09-06 that subtraction is in the chain — `recopseudo.delta_from_amplitude(..., tensor_leakage=...)`, driven by `money_cos2phi_reco.py --subtract-tensor-leakage {kappa,model}`, with `--leakage-scan` printing the per-spot budget — and both routes return the exact-kernel pseudo-measurement to the massless one at the fifth digit at the four sweet spots. The in-situ route first measures across bins, and removes, the bin-independent pedestal that a relative-luminosity error leaves on the fitted constant: uncorrected, $\hat{\kappa}/\kappa_{\text{model}}$ runs 1.46–3.36 over the twelve spots and the subtraction overshoots, calibrated it runs 0.996–1.002, and where the b-sector constant crosses zero and κ is not measurable to the fit's own resolution the correction falls back to the model and says so. The default is off — here as in the generator, so that the chain runs on the massless tensor path it was built on — because what the subtraction leaves is the unmeasured b_3, b_4 band, 0.6 of the leakage: `--leakage-scan` puts it at 0.0014–0.018% of the amplitude at the four bins of the closures below and 0.0003–0.065% over the twelve, so subtracting shrinks a systematic three orders below the bars by a factor 1.6 at the price of re-drawing every pseudo-measurement (plans/08 D10).

4.2 • Kinematic reconstruction and the hadronic final state

The electron method gives Q^2 well ($\delta Q^2/Q^2 = \delta E'/E' \oplus \cot(\theta'/2) \delta\theta'$, measured through the response as an rms of 5.2% at the two $Q^2 = 1.14 \text{ GeV}^2$ spots with the 3 mrad angular placeholder and of 2.9% and 1.6% at the other two, where the η table gives 2 mrad — the formula at the bin centres gives 5.7 / 5.7 / 2.5 / 1.6%) but loses y as $(1 - y)/y$: all four sweet spots sit at $y = 0.010$ – 0.025 (Table 3), where a 1.2% energy resolution gives $\delta y/y = 0.46$ – 1.2 and the bins are not reconstructible. The remedy is HERA's [5,6]: y from the hadronic final state,

$$y_{\Sigma} = \frac{\Sigma}{\Sigma + E'(1 - \cos\theta_e)}, \quad \Sigma = \sum_{h \neq e'} (E_h - p_{z,h}), \quad x = \frac{Q_e^2}{s y_{\Sigma}},$$

with θ_e from the ion direction. Σ is insensitive to particles lost along the ion beam (each contributes $E - p_z = m_T e^{-y}$) and y_{Σ} to collinear initial-state radiation; the mixed x inherits both, because the factor $(1 - z)$ of a radiated photon cancels between Q_e^2 and s — it is the Q^2 label of the bin that migrates under radiation, not x (plans/08 D3 corrected; the migration is bounded in §7). The Jacquet–Blondel and double-angle methods are implemented alongside and checked against hand-derived formulas; the Σ method is the best of the three below $y = 0.05$ and is used throughout.

The hadronic final state is a PYTHIA 8.311 [20] production of 8 M unpolarized $e + p$ and $e + n$ events at the three per-nucleon beam energies ($Q^2 \geq 0.7 \text{ GeV}^2$, lepton radiation off, $\hat{m} \geq 0.5 \text{ GeV}$ so that the library covers the whole x range of the generator window — its 4 GeV default had silently removed $x < 16/s$, 39% of the selected rate, before 2026-08-28), merged in the ratio of their cross sections. It enters as a library: for a pseudo-event at (x, Q^2) a library event of the same cell, passed through the hadron-side response with its noise off, supplies the captured fraction of Σ_h and the p_T ratio, which are applied to the pseudo-event's true $\Sigma_h = 2E_e y + m_N^2/(E_N + p_N)$ — the massless value plus the same target-mass term the library's Σ carries, 4.4 MeV at 99.5 GeV/u — and the 50 MeV noise is added per pseudo-event; y_{Σ} then follows with the

pseudo-event's own reconstructed electron. Table 3 and Figures 1–2 give the result at the mid configuration: 80–92% of the true Σ_h is within the acceptance and above threshold, 7–17% escapes forward beyond $|\eta| = 3.7$ (the target-fragmentation side of a $W = 6\text{--}10$ GeV system; the lithium fragments at $\eta \approx 8$ never enter, each carrying only $m^2/2p$), 1–6% lies below threshold, and 70–85% is captured through the full response. The loss is a scale bias on y_Σ , calibrated as an analysis would calibrate it: the library's own events, smeared and reconstructed like the pseudo-events, give a map of $\langle \Sigma_{\text{reco}} \rangle / \langle \Sigma_{\text{true}} \rangle$ in bins of the reconstructed (x_{mixed}, Q_e^2), and each pseudo-event's sums are divided by the map at its own uncalibrated reconstructed point. (Until 2026-08-28 the map was looked up at the event's *true* cell, which no experiment can do; the measurable version gives purities within 0.06 of it and is the only one used now.) What sets the resolution is not the loss but the noise: with 50 MeV the Σ method gives $\delta y/y = 0.32 / 0.22 / 0.29 / 0.15$ at the four spots against 0.19–0.22 without noise, and at $y = 0.005$ it runs $0.22 \rightarrow 0.33 \rightarrow 0.54 \rightarrow 1.01$ for $0 \rightarrow 25 \rightarrow 50 \rightarrow 100$ MeV. Each configuration selects its own four bins; at the low and top configurations they resolve to $0.39 / 0.23 / 0.24 / 0.11$ and $0.23 / 0.19 / 0.21 / 0.18$, so the $(0.141, 14.3 \text{ GeV}^2)$ bin is best measured at 5×40.8 and the low-x bins improve with $\sqrt{s}$.

SWEET SPOT (X, Q ²)	Y	W [GEV]	E': E', Θ', H	ΔY/Y E' ALONE	Γ _H AS H · Σ- WEIGHTED PARTICLE H (16 / 50 / 84%)	WITHIN ACCEPTANCE, H ≤ 3.7 (TRACKS / Γ / HCAL)	LOST FORWARD	BELOW THRESHOLD	CAPTURED THROUGH THE RESPONSE (MEAN · MEDIAN)	ΔY/Y: Σ, JB, DA
0.028, 1.14	0.010	6.3	9.93 GeV, 107 mrad, −2.93	1.18	1.6 · 1.8 / 2.6 / 3.8	0.80 (0.45 / 0.25 / 0.10)	0.17	0.03	0.70 · 0.73	0.32, 0.32, 0.51
0.011, 1.14	0.025	10.0	9.77 GeV, 108 mrad, −2.92	0.46	0.7 · 0.8 / 1.6 / 2.9	0.87 (0.51 / 0.26 / 0.10)	0.08	0.06	0.74 · 0.78	0.22, 0.22, 0.40
0.071, 3.13	0.011	6.5	9.97 GeV, 177 mrad, −2.42	1.07	2.1 · 1.9 / 2.6 / 3.7	0.83 (0.47 / 0.25 / 0.10)	0.15	0.02	0.74 · 0.77	0.29, 0.29, 0.44
0.141, 14.3	0.025	9.4	10.10 GeV, 378 mrad, −1.65	0.46	2.0 · 1.7 / 2.2 / 2.9	0.92 (0.55 / 0.27 / 0.09)	0.07	0.01	0.85 · 0.87	0.15, 0.15, 0.21

Table 3 — the four sweet spots of Report 1 at $e(10) \times {}^6\text{Li}(99.5)$ and their hadronic final state (PYTHIA 8, 4 M e + p and e + n events; `hfs_acceptance.py`, `hfs_resolution.py`). The acceptance columns are geometry and thresholds only, in the detector frame; "captured through the response" includes efficiencies, the track turn-on and the mass hypotheses, noise off, and equals the median y_Σ/y ; the resolutions are 68% half-widths through the full response with 50 MeV of noise. Differences of 0.01 between this table and the figures are Monte-Carlo noise.

Hadronic-method resolution with a hadronic final state — $e(10) \times 6\text{Li}(99.5/u)$; pythia8, 4000000 events
 tracker $|\eta| \leq 3.5$ $p_T > 0.20$ GeV eff 0.95; cal $|\eta| \leq 3.7$, photon $E > 0.10$, n-had $E > 0.5$ eff 0.90; noise 50 MeV




Figure 1 — hadronic-method resolution ($e10 \times {}^6\text{Li}99.5$, PYTHIA 8, `hfs_resolution.py`). **(a)** $\delta y/y$ (68% half-width; dotted: rms) of the Σ , Jacquet-Blondel and double-angle methods at $Q^2 < 3 \text{ GeV}^2$ with 50 MeV of noise, against the electron method and the 25% stand-in of the earlier chain. **(b)** Captured fraction of Σ by detector object, rate-weighted, no noise. **(c)** The four sweet spots. **(d)** The Σ -method resolution for 0, 25, 50 and 100 MeV of noise: the floor decides the $y \approx 0.01$ bins.

Hadronic-final-state acceptance at the four sweet spots — $e(10) \times 6\text{Li}(99.5/u)$, PYTHIA 8; tracker $|\eta| \leq 3.5$, $p_T > 0.2 \text{ GeV}$; photons > 0.1 , neutral hadrons $> 0.5 \text{ GeV}$




Figure 2 — where the hadronic $E - p_z$ sum goes ($e10 \times {}^6\text{Li}99.5$, PYTHIA 8, `hfs_acceptance.py`). **(a)** The true Σ_h by fate — captured as tracks, photons and HCal objects; lost forward beyond the calorimeters, below threshold, or backward — for a calorimeter reach of $|\eta| \leq 3.7$ (full) and the ePIC nominal 4.0 (light). **(b)** The Σ -method $\delta y/y$ through the full response for both reaches: the extra reach moves the bias, not the resolution.

4.3 · The spin-state-sorted estimator

A single-fill fit of $N(\varphi')$ measures the detector's own $\cos 2\varphi'$ acceptance harmonic divided by P_{ZZ} on top of the physics: a 3% harmonic aligned with the vertical axis — the symmetry axis of the crossing-angle plane, of module boundaries and of the Roman-Pot cutout — fakes 0.05, five to eleven times the projected amplitudes (Figure 4d). The measurement therefore sorts the events of each reconstructed bin by spin state f , histograms $N_f(\varphi')$, and forms the bin-wise ratio

$$R(\phi') = \frac{\sum_f w_f N_f(\phi')}{\sum_f N_f(\phi')}, \quad w_f = P_f - \bar{P}, \quad \bar{P} = \frac{\sum_f L_f P_f}{\sum_f L_f}, \quad R = \frac{\sigma_P^2 T(\phi')}{1 + \bar{P} T(\phi')}, \quad T = \kappa + \hat{A} \cos 2\phi', \quad \sigma_P^2 = \frac{\sum_f L_f (P_f - \bar{P})^2}{\sum_f L_f},$$

which is exact for yields $N_f = L_f \epsilon(\phi') \sigma [1 + P_f T]$ because $\sum_f L_f w_f = 0$: any efficiency common to the spin states cancels bin by bin, and a relative-luminosity error enters only through $\bar{P}$, as a ϕ' -independent constant plus a second-order rescaling. R is inverted for T bin by bin (with the exact Jacobian in the error propagation) and T fitted by weighted least squares to $\kappa + \hat{A} \cos 2\phi' + \hat{A}_s \sin 2\phi'$, the finite-bin dilution $\sin w/w$ divided out; the bin variances use the expected per-state counts so that a bin populated by one state alone keeps a finite weight. Two states of equal luminosity give $\delta\hat{A} = 2\sqrt{(2/N)/(P_+ - P_0)}$: the $m = \pm 1$ -rich state at $P_+ = +0.6$ alternated with the $m = 0$ -rich state at $P_0 = -1.2$ is $1.5\times$ more precise than a single $P_{zz} = 0.6$ fill with the whole luminosity, and an unpolarized reference state would cost a factor two instead. κ is the transverse tensor rate shift of the b_1 sector, measured as a by-product; its sign is that of $-b_1 P_2(\cos \theta_S)$ in the tensor convention this simulation adopted on 2026-08-29 ($A_{zz} = -(2/3) b_1/F_1$, Cosyn et al. [12] and HERMES), so a transverse fill with a positive b_1 returns a positive κ . The cancellation holds only for an acceptance *common* to the states: a difference $\delta\epsilon_2$ of the $\cos 2\phi'$ harmonic between them enters as $\sum_f L_f (P_f - \bar{P}) \epsilon_f / \sigma_P^2 = \delta\epsilon_2 / (P_+ - P_0)$ at any luminosity split, so the states must alternate bunch by bunch within a fill, or the harmonic be stable to 10^{-4} between fills (§6). The same estimator in two azimuths serves the coherent channel (§4.5).

R itself is unbiased at any occupancy — conditional on a bin's total the per-state split is multinomial and R is linear in it — but the *inversion* is not. $\hat{T} = R(1 + u)/(\sigma_P^2 - \bar{P}R)$ is strongly curved over the range $|R| \leq \max_f |P_f - \bar{P}|$ that R explores when a bin holds a handful of counts (a bin with one count has $R = \pm 0.9$ exactly): its second derivative $2(1 + u)\sigma_P^2 \bar{P}/(\sigma_P^2 - \bar{P}R)^3$ carries the sign of $\bar{P}$, so the flip plan's $P_0 = -2P_+ - \bar{P} = -0.3$ — makes the inversion concave and each bin carry a negative Jensen offset $(1 + u)\bar{P}/(\sigma_P^2 v)$, while the data-driven $1/\text{var}$ weights add an opposite-sign term of the same order. The inclusive ϕ' bins carry $\geq 10^4$ counts and are immune; the sparse coherent $|t|$ bins are not, and there the offset is common to every α at fixed β (the acceptance is flat in α and modulated $\times 25$ in β by the pot cutout), so it is orthogonal to $\langle \cos 2\alpha \rangle$ and lands on a_t and the constant. The alternative is to keep the per-bin efficiency as a free parameter and profile it out. Because ϵ_b enters as a linear scale the profile is closed-form, $\epsilon_b = n_b/(1 + u_b + \bar{P}T_b)$, and substituting it back leaves *exactly* the conditional multinomial given the bin totals, $p_{f,b} = L_f(1 + u_b + P_f T_b)/(1 + u_b + \bar{P}T_b)$. Profiling as many nuisance parameters as there are bins would ordinarily raise the incidental-parameter objection; it does not here, because the conditional score has zero mean bin by bin at any count — $\sum_f p_{f,b}(P_f/q_f - \bar{P}/D)$ vanishes identically — so no $1/v$ term survives. Newton from $T = 0$ converges in six iterations under a step-halving guard that keeps every populated bin's density positive — and a step that guard has halved to nothing is not read as convergence: the gradient at the solution is tested against the error bars it would move, so a fit driven onto the boundary is named rather than returned; the covariance is the conditional Fisher information at the observed totals; empty bins contribute nothing and single-state bins contribute a finite $n \log p$ term, which removes the infinite-weight pathology structurally rather than by a variance floor. It is the same model, the same basis and the same assumptions as the ratio — in particular the same blindness to a state-dependent acceptance, measured identical to within 2% of the statistical error — and on exact counts it returns the injected coefficients to machine precision with errors that agree with the ratio's to 0.24%, so no error bar in this report moves. It is available throughout as an option, the ratio remains the default, and §5.2 measures what it buys.

4.4 • Response, calibration and unfolding

The inclusive response is importance-sampled: every accepted cell of a 60×45 log-log grid receives 400 pseudo-events, log-uniform in the cell, with weight σ_c/n_c (n_c the draws inside the generator acceptance, so the cell rate is conserved; the amplitude is taken at the cell centre, worth $\leq 0.2\%$ at the sweet spots), and each passes through the chain of Schematic 1. For a reconstructed bin the response gives its purity (the fraction of its rate that originates in the true bin), its efficiency (the fraction of the true-bin rate reconstructed in it after ϵ_{eID}), and D , the rate- and ϵ_{eID} -weighted, ϕ' -diluted amplitude of the events landing in it over the amplitude of the true bin. The expected ϕ' counts per spin state are exact bin integrals of $\text{den}_f + P_f \hat{A} \cos 2\phi'$ over those events, with the ϕ' efficiency and the luminosity shares of the run plan; they are Poisson-fluctuated

at 10 and 100 fb⁻¹/u and the fit is compared with the reconstructed-bin truth, which is what an unbiased estimator on this sample must return.

The amplitude is exactly linear in Δ — Δ enters the master formula in one place — so the response is a linear operator on Δ , $\text{fold}(\Delta) = \langle (\partial A / \partial \Delta) \Delta(x_{\text{true}}, Q^2_{\text{true}}) \rangle$ over the events of the bin, and the conversion $\Delta = \hat{A} K$ with $K = \Delta(x_c, Q^2_c) / \text{fold}(\Delta)$ is the bin-centering with migration included. Evaluating K with the shape that generated the pseudo-data is circular and reproduces the injected curve by construction; evaluating it with a wrong prior biases Δ by up to 20–29% over the plotted bins of a Q^2 slice, and by (–4.2, +8.0, –5.6, +4.9)% (model B) or (+8.6, +2.5, +7.3, +3.1)% (flat toy) at the four spots. Fitting a two-parameter tilt of the prior *through* the response per Q^2 slice on a bounded grid — the data choose the shape, the box guards the conditioning — reduces the residual to at most 4.0 / 5.7 / 5.8% (model B) and 4.0 / 4.3 / 3.6% (toy) over the plotted bins of the three slices, and to (–1.5, –1.8, –0.3, +0.5)% and (+4.6, +2.9, +1.8, –0.9)% at the spots. The Δ error bar then carries the statistics, the shape-fit error, the response Monte Carlo error (0.1–1.0% per plotted bin) and the spread over the priors the tilt cannot absorb, the last dominating in the best-measured bins.

4.5 • The coherent channel

The intact recoil keeps the beam rigidity ($A/Z = 2$, $x_p \leq 0.01$), so the pots resolve only its transverse displacement: the reconstruction returns the angle pair (θ_x, θ_y) at the interaction point relative to the beam axis — crossing angle and reference orbit drop out by construction — smeared by the bunch's own divergence, and $p_T = A p_u |\theta|$, $\phi_t = \text{atan2}(\theta_y, \theta_x)$, $|t| = p_T^2$ with x_L set to one ($t_{\min} = M^2 x_p^2 / (1 - x_p) = 3 \times 10^{-3} \text{ GeV}^2$ at $x_p = 0.01$ is a known bias once x_p is measured from M_X). A recoil is tagged when it clears, per axis, the larger of the 10σ beam envelope and the pot aperture. At the Yellow Report optics with the measured aperture the tagged fraction is 9×10^{-10} , 6×10^{-27} and 7×10^{-14} at the three configurations (Report 3 §5; Figure 3c marks the envelope and the aperture separately): the coherent channel exists only under the tagging optics of Report 1 §6.1, in which the horizontal β^* is raised to the optimum of acceptance $\times$ luminosity, the vertical plane stays at high acceptance, and the pots follow the envelope in both planes — a cutout of $0.36 \times 3.8 \text{ mrad}$ ($0.09 \times 0.93 \text{ GeV}$) at 5×40.8 with $\sigma_0 = 36 / 380 \text{ } \mu\text{rad}$ (the horizontal including the beam's momentum spread through the pot dispersion), giving a 37% tag at 1/6.8 of the high-acceptance luminosity. The closure below is at that point.

With the recoil measured there are two physical azimuths besides the spin: $\alpha = \phi_e - \phi_S$ and $\beta = \phi_t - \phi_S$, and the spin-sorted yield is modelled as

$$N_f(\alpha, \beta) \propto \varepsilon(\alpha, \beta) [1 + u_1 \cos(\alpha - \beta) + u_2 \cos 2(\alpha - \beta) + P_f(\kappa + a_e \cos 2\alpha + a_t(t) \cos 2\beta + a_m \cos(\alpha + \beta))],$$

where a_t is the deformation term of the polarized-deuteron anchor [9] ($\propto |t|$ at low $|t|$, coefficient $2a_2$ in the anchor's convention, predicted to flip sign for ${}^6\text{Li}$), a_e the flat double-helicity-flip term that needs the lepton plane, a_m a nuisance, and u_1, u_2 the L–T and T–T' interference of unpolarized diffractive DIS [13,14]. The two mechanisms are separated by the azimuth they modulate and by their t dependence. The cutout is far from circular, so its acceptance carries a large $\langle \cos 2\beta \rangle$ of its own — in the true recoil azimuth –0.49 at the tagging cutout and –0.95 at the measured aperture, the divergence smearing diluting the first by $\langle \cos 2(\beta_{\text{reco}} - \beta_{\text{true}}) \rangle = 0.64$ to the –0.31 money plot 6R annotates — and no single-fill fit survives; the spin-state ratio of §4.3 in (α, β) cancels it bin by bin, and its acceptance-profiled likelihood twin does the same without the low-count bias that the sparsest $|t|$ bins carry (Table 5). Because the cutout correlates β with $|t|$ inside a reconstructed $|t|$ bin, the fit uses an acceptance-weighted basis from the response — the in-bin means of $\cos \beta$, $\sin \beta$, $\cos 2\beta$, $\sin 2\beta$ of the true β per reconstructed β bin, and the template $\langle (t/t_{\text{ref}}) \cos 2\beta \rangle$ of the assumed linear shape — so that the coefficients refer to the unsmeared truth and a_t is quoted at t_{ref} , the rate-weighted mean true $|t|$ of the bin; this is one coefficient per bin, the t dependence of a_t coming from the bins together. There are seven of them, on the edges 0.017, 0.028, 0.039, 0.05, 0.08, 0.12, 0.17 and 0.25 GeV² of `recopseudo.T_EDGES_PUBLISHED`, adopted 2026-08-28 (§7); its lower edge is a choice of the analysis and not a limit of the apparatus, the cutout's own floor ($A p_u \varepsilon_x$)² sitting at 0.0079, 0.0132 and 0.0094 GeV² at the three configurations. Where the cutout empties part of the β circle the design absorbs it: sixteen to twenty of the twenty-four β bins are populated below 0.05 GeV² and twenty to twenty-four above it, all twenty-four in the top bin, and the weighted design keeps rank 7/7 at condition numbers 1.85 to 10.87 at every configuration. What the window costs is model rather than conditioning — the digitized deuteron anchor

coherent.MANTYSAARI_A2_DEUTERON begins at $|t| = 0.05 \text{ GeV}^2$, so wherever t_{ref} falls below that the injected a_t is the linear deformation shape extrapolated below its lowest anchor point. Counted per bin that is four of the seven and not three, because the resolution puts the t_{ref} of the 0.05–0.08 bin at 0.046 GeV^2 alongside the added bins' 0.023, 0.030 and 0.036; at 18×137.5 , where the tighter envelope moves every t_{ref} up, the fourth is marginal — 0.0498 GeV^2 against the anchor's lowest point at 0.050 — and the other three are 0.0233, 0.0310 and 0.0383; the retired window's own lowest bin is that fourth one, so the wider window deepens the extrapolation rather than introducing it. The three parity-forbidden partners $\sin 2\alpha$, $\sin 2\beta$ and $\sin(\alpha + \beta)$ are fitted alongside as a null test: a spin-axis error δ gives $\tan 2\delta$ in both azimuths, a pot roll δ_t gives $\tan 2\delta_t$ in β alone. α is integrated analytically per event; the coherent pseudo-experiment has no electron side, and x_p is sampled but never reconstructed (the M_X method would provide it).

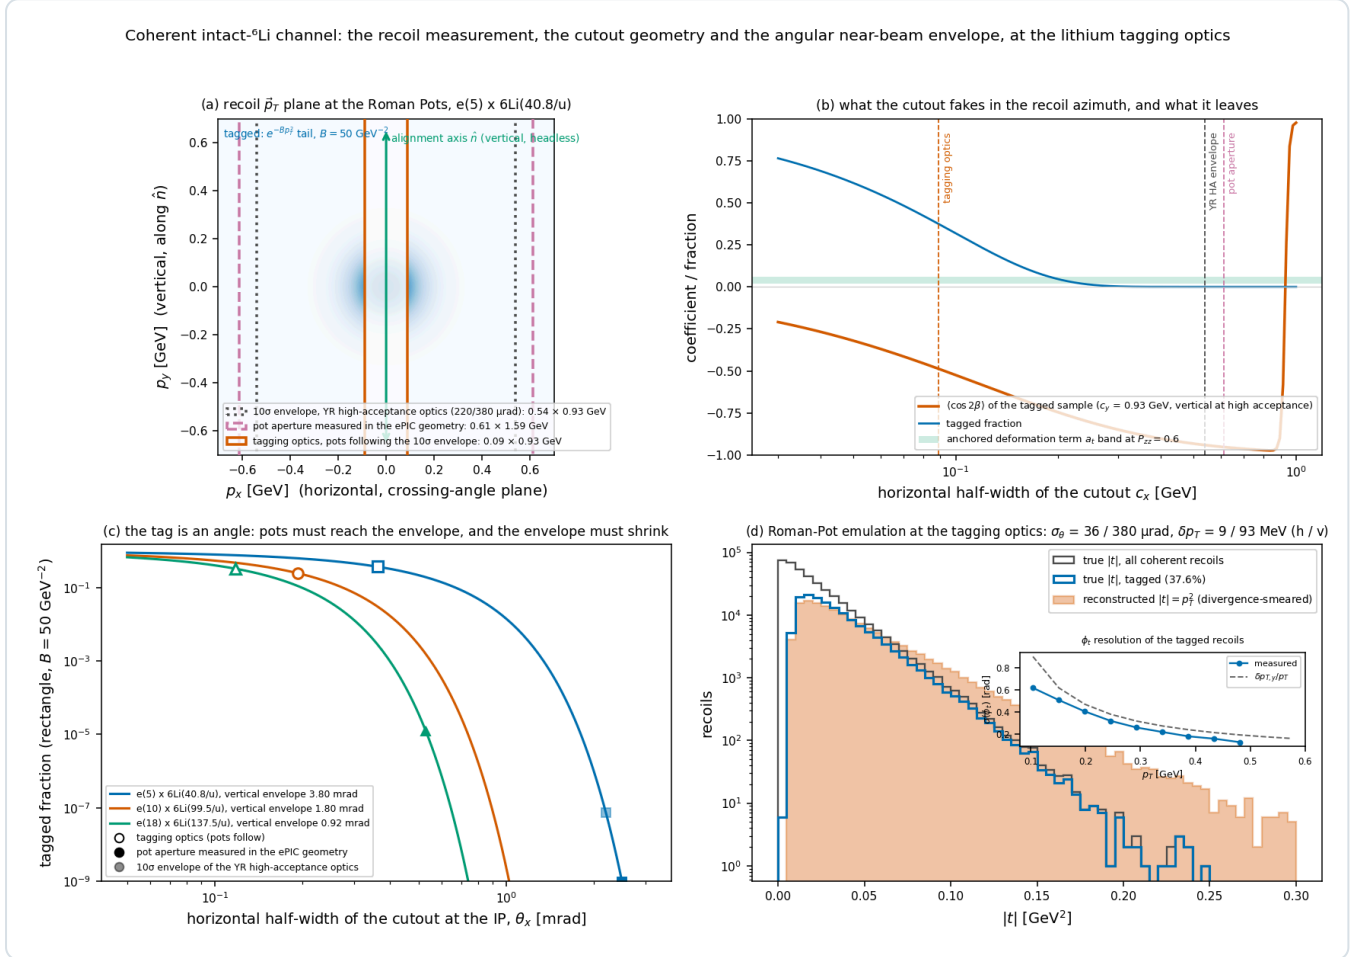


Figure 3 — the coherent measurement quantified (5×40.8 , `reco_chain_figures.py`). **(a)** The recoil p_T plane with the exponential coherent spectrum and the three cutouts: the Yellow Report high-acceptance 10σ envelope, the pot aperture measured in the ePIC geometry, and the tagging optics with pots following the envelope. **(b)** $(\cos 2\beta)$ of the tagged sample about the vertical spin axis and the tagged fraction versus the horizontal half-width of the cutout, against the anchored deformation band. **(c)** Tagged fraction versus the horizontal half-angle at the interaction point for the three configurations, with the three cutouts marked; the markers below 10^{-9} are off the axis, and the aperture marker at 18×137.5 lies inside the envelope that binds there. **(d)** The Roman-Pot emulation at the tagging optics: true, tagged and reconstructed $|t|$, and (inset) the ϕ_t resolution, dominated by the 93 MeV vertical divergence. The tagged fraction the panel prints, 37.6%, is per recoil above t_{min} ; Table 5's 36% is the same Monte Carlo and the same cutout counted per recoil produced, the 3.3% between them being the draws the t_{min} cut removes.

4.6 · What the analysis reads

The paths from counts to $\hat{A}$ and to a_t , a_e were traced assignment by assignment in the review (Appendix A). They consume: the spin state of each bunch crossing, the tensor polarization and luminosity share of each state (the run-plan bookkeeping, known to an experiment from the source pattern, the polarimeter and the luminosity monitor); the reconstructed electron four-vector; the hadronic sums and the calibration map at the reconstructed point; the Roman-Pot angle pair after the cut; and the simulation-derived corrections — purity, efficiency, K , the β basis, t_{ref} — in reconstructed bins. Generator truth is

read to generate the pseudo-data, to build those corrections, and to plot the closure references; the analysis assumes the nuisances u_1 , u_2 and the luminosity shares (as fit arguments, so that a wrong assumption can be exercised) and takes P_{zz} as exactly known (its scale is a normalisation, §6). The bins shown in the figures are selected on their expected count and statistical error alone.

5 • Closure at the reconstructed level

5.1 • Inclusive: money plots 5R and 7R

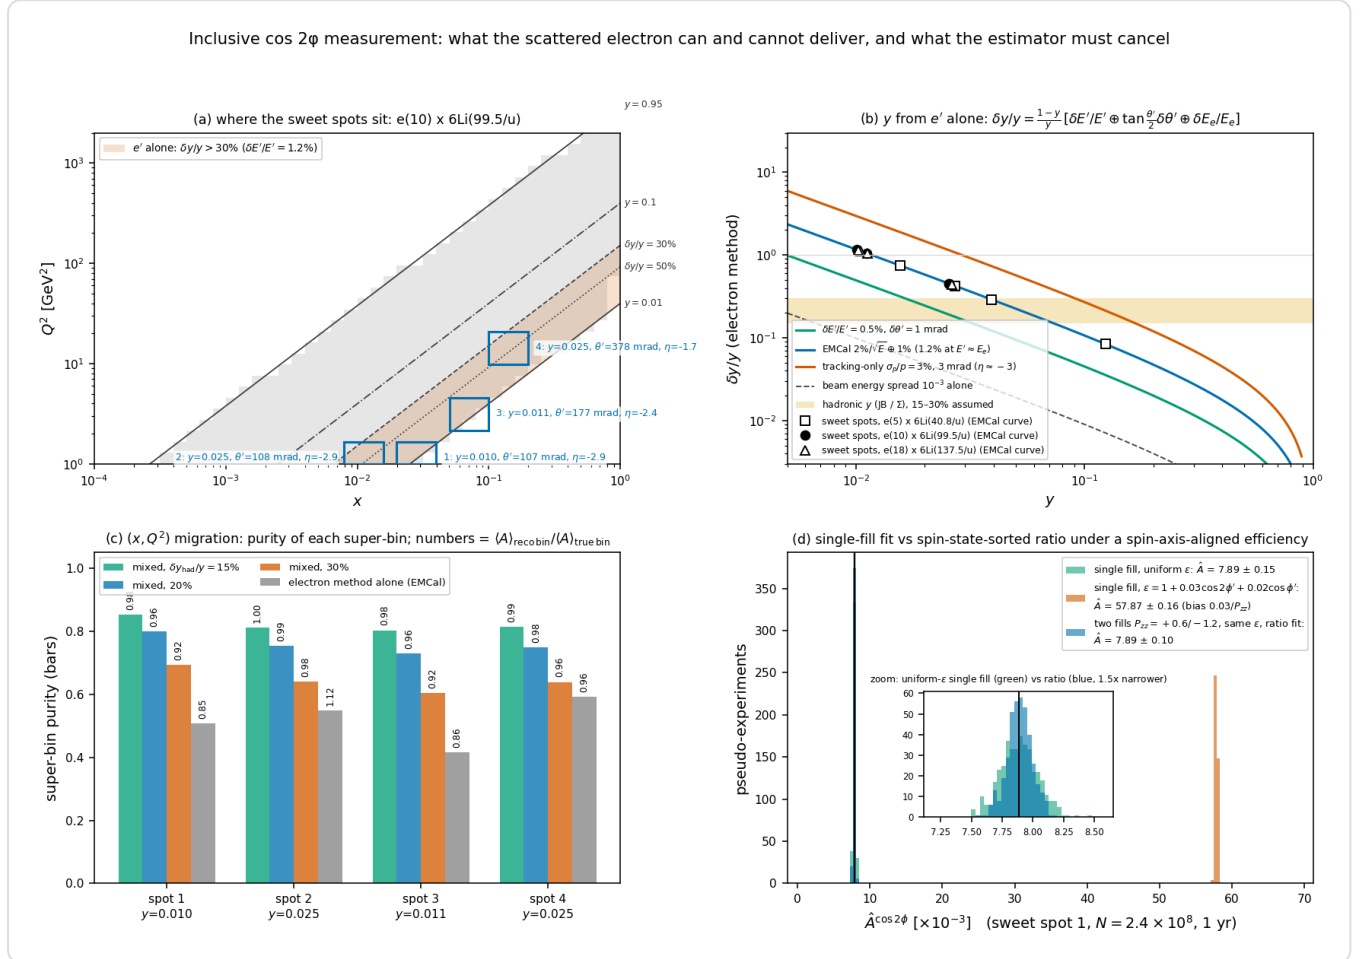



Figure 4 — the inclusive measurement quantified ($e10 \times {}^6\text{Li}99.5$, `reco_chain_figures.py`). **(a)** The analysis acceptance with iso- y lines and the four sweet-spot super-bins; the shaded band is where the electron alone gives $\delta y/y > 30\%$ at a 1.2% energy resolution. **(b)** The electron-method $\delta y/y$ versus y for calorimeter, tracking and beam-energy-spread inputs, with the sweet spots of the three configurations and the 15–30% hadronic band. **(c)** Purity of each super-bin for three Gaussian hadronic- y resolutions and for the electron method alone; the numbers are the reconstructed-bin over true-bin amplitude; they are migration purities — the electron smearing and the hadronic y alone, without the reconstructed-level analysis window or ϵ_{elD} — and are not on the footing of the purities of Table 4, which are through the full selection. **(d)** Four hundred pseudo-experiments of the spot-1 super-bin at one year, generated at its truth-level count $N = 2.4 \times 10^8$ (Table 4's 1.6×10^8 is the reconstructed one, so the errors here are 0.10 rather than 0.12×10^{-3}): the single-fill fit under a uniform efficiency (green), the same under $\epsilon = 1 + 0.03 \cos 2\phi' + 0.02 \cos \phi'$ (vermillion, biased by $0.03/P_{zz}$), and the two-state ratio under that efficiency (blue, unbiased and $1.5 \times$ narrower).

SWEET SPOT (RECO BIN)	25% GAUSSIAN HADRONIC-Y STAND-IN				PYTHIA 8 FINAL STATE, Σ METHOD, 50 MEV NOISE, CALIBRATED					$\Delta\hat{A}$ (10 YR) [10 ⁻⁵]	ΔA , REPORT 1 [10 ⁻⁴]
	PURITY	EFF.	D	$\hat{A} \pm \Delta\hat{A}$ (1 YR) [10 ⁻³]	N (1 YR)	PURITY	EFF.	D	$\hat{A} \pm \Delta\hat{A}$ VS RECO-BIN TRUTH [10 ⁻³]		
1: x = 0.028, Q ² = 1.14	0.66	0.42	0.92	6.88 ± 0.12 (7.00)	1.6×10 ⁸	0.56	0.31	1.00	7.38 ± 0.12 (7.61)	3.9	1.7


2: $x =$ 0.011, $Q^2 =$ 1.14	0.63	0.59	0.99	$4.30 \pm$ 0.09 (4.20)	3.0×10^8	0.59	0.56	1.06	4.61 ± 0.09 (4.50)	2.9	1.4
3: $x =$ 0.071, $Q^2 =$ 3.13	0.69	0.38	0.90	$8.71 \pm$ 0.16 (9.07)	8.8×10^7	0.64	0.32	0.95	9.78 ± 0.17 (9.56)	5.4	2.7
4: $x =$ 0.141, $Q^2 =$ 14.3	0.69	0.66	0.96	$9.07 \pm$ 0.29 (9.23)	2.9×10^7	0.75	0.71	0.98	9.78 ± 0.30 (9.41)	9.4	4.5

Table 4 — money plot 5R: the four sweet-spot super-bins in reconstructed (x , Q^2) at $e(10) \times {}^6\text{Li}(99.5)$, $P_{zz} = 0.6$, model A, mixed method, spin-state ratio of the $+0.6 / -1.2$ states with the acceptance harmonic and the 10^{-3} luminosity offset on, statistical errors only. The reconstructed-bin truth is in parentheses. With the PYTHIA final state uncalibrated the purities would be $0.42 / 0.53 / 0.49 / 0.68$ and $D = 0.79 / 0.84 / 0.83 / 0.95$; the errors are the same to within 7%. The last column is the single-state, truth-binned error of Report 1.

Three observations. The fit is unbiased: the eight pulls against the reconstructed-bin truth lie within 2.3 standard errors with the acceptance harmonic and the luminosity offset on, and the $\sin 2\phi'$ coefficients are consistent with zero. The errors are 0.60–0.72 of the single-state truth-level values of Report 1 over the two hadronic- y models, and 0.60–0.67 with the Gaussian stand-in alone, essentially the $1.5\times$ gain of the $m = 0$ -rich state on its own: the reconstruction efficiency of 0.3–0.7 is measured against the truth bin before the analysis cuts and costs the projection no statistics, migration into the reconstructed bin leaving it with $0.97\text{--}1.25\times$ the truth-level rate — presuming the source delivers that state at $|P_{zz}| = 1.2$ with the same purity as the ± 1 state at 0.6. And the real final state costs purity rather than statistics: the $y \approx 0.01$ spots lose about 60% of their events to the reconstructed $y \geq 0.01$ cut and a further third of the remainder to the Σ resolution, which is why more of the answer passes through the bin-centering factor there. Unfolded to Δ (money plot 7R), the best reconstructed bins reach statistical errors $\delta\Delta = 2.3 \times 10^{-3}$ at $Q^2 = 1.14 \text{ GeV}^2$, 1.1×10^{-3} at 3.13 GeV^2 and 0.3×10^{-3} at 14.3 GeV^2 in one year ($1.7 / 1.7 / 6.9\%$ of Δ ; purities 0.44–0.56; $2.4 / 1.2 / 0.3 \times 10^{-3}$ at 0.54–0.57 with the stand-in). With the bin-centering factor fitted through the response the full one-year bars at those bins are $4.2 / 2.6 / 7.0\%$ and the ten-year ones $3.8 / 2.2 / 2.6\%$, the prior spread dominating at the two lower Q^2 — $3.7 / 2.1 / 1.3\%$ of the ten-year systematic, against a one-year systematic of $3.9 / 2.0 / 1.4\%$ in all, the factor being refitted on each luminosity's own draw; at $Q^2 = 1.14$ and 3.13 GeV^2 the ten-year data reject the flat-toy shape ($\chi^2/\text{ndf} = 31/9$ and $42/10$ against $6/9$ and $16/10$ for the injected one), at 14.3 GeV^2 no member of the family is distinguishable.



Money plot 5R — inclusive gluonometry at the reconstructed level ($e10 \times 6\text{Li}99.5$, $P_{zz} = 0.6$, model A, PYTHIA 8 final state with the Σ method and the calibrated hadronic scale, EMCal energy, track angles, $\epsilon_{\text{eID}}(\eta)$, spin-state ratio of the $+0.6/-1.2$ states with the acceptance harmonic and luminosity offset on). Left: the acceptance-free modulation $\pi(\phi') - \hat{\kappa}$ per unit P_{zz} with the reconstructed-bin truth (blue) and the fit (vermillion). Right: the amplitude in reconstructed x bins along the $Q^2 = 1.14$ GeV² slice against the reconstructed-bin truth (blue ticks) and the model curves.



Money plot 7R — $\Delta(x, Q^2)$ from reconstructed bins. $x\Delta$ points at the three sweet-spot Q^2 slices from the same pseudo-measurements, converted with the bin-centering factor fitted through the response (§4.4), so that the error bars carry the statistics, the shape fit, the response Monte Carlo and the prior spread; one-year (open) and ten-year (filled) draws are independent.

5.2 · Coherent: money plot 6R


CONFIGURATION	RECO T BIN [GeV ²]	N _{TAG} , 1 YR	A _T (T _{REF}) INJECTED	FIT, 1 YR	ENSEMBLE MEAN, RATIO (20 EXP.)	ENSEMBLE MEAN, LIKELIHOOD (SAME 20 DRAWS)	FIT, 10 YR	A _E FIT, 1 YR (INJ. 0.010)	ENSEMBLE MEAN	A _E FIT, 10 YR
5 × 40.8, tag 36%, N _{tag} = 2.3 × 10 ⁶ /yr	0.017– 0.028	5.5 × 10 ⁵	0.046	0.0418 ± 0.0063	0.0453 (σ 0.0054)	0.0451 (σ 0.0054)	0.0441 ± 0.0020	0.0099 ± 0.0022	0.0105 (σ 0.0018)	0.0090 ± 0.0007
	0.028– 0.039	4.1 × 10 ⁵	0.060	0.0593 ± 0.0052	0.0591 (σ 0.0044)	0.0586 (σ 0.0043)	0.0589 ± 0.0016	0.0125 ± 0.0026	0.0093 (σ 0.0027)	0.0125 ± 0.0008

18 × 137.5, tag 32%, N _{tag} = 6.0×10 ⁶ /yr	0.039– 0.05	2.9×10 ⁵	0.073	0.0714 ± 0.0052	0.0709 (σ 0.0061)	0.0709 (σ 0.0061)	0.0738 ± 0.0016	0.0087 ± 0.0031	0.0089 (σ 0.0033)	0.0091 ± 0.0010
	0.05– 0.08	4.1×10 ⁵	0.092	0.0973 ± 0.0037	0.0913 (σ 0.0039)	0.0916 (σ 0.0039)	0.0922 ± 0.0012	0.0111 ± 0.0026	0.0096 (σ 0.0029)	0.0100 ± 0.0008
	0.08– 0.12	1.8×10 ⁵	0.121	0.1209 ± 0.0050	0.1203 (σ 0.0042)	0.1215 (σ 0.0042)	0.1207 ± 0.0016	0.0126 ± 0.0040	0.0107 (σ 0.0039)	0.0075 ± 0.0013
	0.12– 0.17	5.3×10 ⁴	0.150	0.1535 ± 0.0089	0.1460 (σ 0.0070)	0.1514 (σ 0.0066)	0.1510 ± 0.0028	−0.0052 ± 0.0072	0.0084 (σ 0.0082)	0.0114 ± 0.0023
	0.17– 0.25	1.3×10 ⁴	0.183	0.1166 ± 0.0186	0.1211 (σ 0.0220)	0.1814 (σ 0.0233)	0.1799 ± 0.0060	0.0207 ± 0.0141	0.0117 (σ 0.0182)	0.0093 ± 0.0045
	0.017– 0.028	1.5×10 ⁶	0.047	0.0434 ± 0.0039	0.0463 (σ 0.0041)	0.0461 (σ 0.0042)	0.0472 ± 0.0012	0.0097 ± 0.0013	0.0101 (σ 0.0013)	0.0097 ± 0.0004
	0.028– 0.039	1.2×10 ⁶	0.062	0.0637 ± 0.0031	0.0610 (σ 0.0037)	0.0606 (σ 0.0037)	0.0644 ± 0.0010	0.0112 ± 0.0016	0.0102 (σ 0.0013)	0.0094 ± 0.0005
	0.039– 0.05	8.0×10 ⁵	0.077	0.0742 ± 0.0032	0.0771 (σ 0.0031)	0.0770 (σ 0.0031)	0.0760 ± 0.0010	0.0127 ± 0.0019	0.0094 (σ 0.0021)	0.0099 ± 0.0006
	0.05– 0.08	1.1×10 ⁶	0.100	0.0985 ± 0.0023	0.1001 (σ 0.0024)	0.1002 (σ 0.0024)	0.0987 ± 0.0007	0.0121 ± 0.0016	0.0094 (σ 0.0019)	0.0107 ± 0.0005
	0.08– 0.12	4.1×10 ⁵	0.137	0.1403 ± 0.0033	0.1368 (σ 0.0027)	0.1372 (σ 0.0027)	0.1369 ± 0.0011	0.0126 ± 0.0026	0.0090 (σ 0.0026)	0.0109 ± 0.0008
	0.12– 0.17	1.0×10 ⁵	0.179	0.1666 ± 0.0063	0.1776 (σ 0.0066)	0.1787 (σ 0.0062)	0.1762 ± 0.0020	0.0062 ± 0.0052	0.0094 (σ 0.0048)	0.0078 ± 0.0016
	0.17– 0.25	1.9×10 ⁴	0.228	0.1902 ± 0.0148	0.2060 (σ 0.0184)	0.2289 (σ 0.0178)	0.2221 ± 0.0047	−0.0057 ± 0.0116	0.0115 (σ 0.0115)	0.0051 ± 0.0037

Table 5 — money plot 6R: the two-azimuth template fit per reconstructed |t| bin at the tagging optics (coefficients per unit P_{ZZ}; ε_{B0} = −0.08, B = 50 GeV^{−2}; u₁ = 0.05, u₂ = 0.02; luminosity at the tagging optics, 1/6.8 and 1/9.5 of the high-acceptance value; 6 × 10⁶ response recoils; `money_cos2phi_coherent_reco.py --optics tagging`). The “fit” columns are one pseudo-experiment; the ensemble columns are the mean and spread of twenty one-year pseudo-experiments on the same response, and say whether the estimator is biased. In the five bins at 5 × 40.8 that hold more than 10⁵ tagged recoils the ratio is not, at twenty draws: the means sit within two standard errors of the injected values for a_t and a_e alike, and no single-experiment a_e among them departs from the injected 0.010 by more than one standard error — the largest, 0.0125 ± 0.0026 in 0.028–0.039, is a fluctuation of that one draw whose twenty-experiment mean is 0.0093. Below that count it is biased, and the two upper bins of 5 × 40.8 measure the defect at one year: 1900, 1422, 1005, 1431, 607, 184 and 45 counts per (α, β) cell give ratio means differing from the injected values by −2.0, −1.3, −2.5, −0.3, −0.2, −2.8 and −33.9%, while the same bins close at ten years. The likelihood column is the fix (§4.3) on the *same* twenty Poisson draws: 0.0451, 0.0586, 0.0709, 0.0916, 0.1215, 0.1514 and 0.1814 against 0.0462, 0.0599, 0.0727, 0.0916, 0.1205, 0.1502 and 0.1833 injected, and 0.0461, 0.0606, 0.0770, 0.1002, 0.1372, 0.1787 and 0.2289 against 0.0465, 0.0619, 0.0766, 0.0996, 0.1371, 0.1791 and 0.2283 at 18 × 137.5. Two hundred experiments turn that into a measurement: the ratio's means are −0.6, +1.3, +0.6, −0.3, −1.5, −2.5 and −34.4% from the injected values with pulls of the mean −0.6, +2.1, +1.2, −1.2, −5.4, −5.8 and −47.2, the likelihood's −0.9, +0.7, +0.6, 0.0, −0.5, +1.3 and −0.6% with pulls −1.0, +1.0, +1.1, −0.1, −1.8, +3.1 and −0.9 — six of the seven inside two, and the seventh, the +1.3% of the 0.12–0.17 bin, is the response Monte Carlo's own floor at 6 × 10⁶ recoils rather than a residual of the estimator (plans/08 A12). The error bars become honest with it: the likelihood's spreads are 0.0060, 0.0051, 0.0053, 0.0035, 0.0047, 0.0089 and 0.0186 against quoted errors 0.0063, 0.0052, 0.0052, 0.0037, 0.0050, 0.0089 and 0.0191, where the ratio's quoted error in the sparsest bin is 2% below its own spread of 0.0189 at 45 counts per cell. a_e is unbiased under both estimators, as the α-flat acceptance requires. Coarser binning only attenuates the defect: over the same two hundred experiments the sparsest bin's −34.4% at the default 12 × 24 (α, β) grid becomes −12.8% at 8 × 16 and −6.9% at 6 × 12, at a

cost of $0.0141 \rightarrow 0.0150 \rightarrow 0.0162$ on that bin's δa_e and of $0.00121 \rightarrow 0.00129 \rightarrow 0.00140$ on the combined one, 7 and 16% through the wider-bin dilutions, and a 4×8 grid is rank-deficient at any count because $\langle \cos 2\alpha \rangle$ vanishes identically over four α bins; it is a cross-check, not the remedy. The ratio remains the default, so every number in the "fit" and ratio-ensemble columns is the published one. The fitted sine coefficients, not shown, are consistent with zero throughout and resolve a pot roll to between 15 and 51 mrad per year across the seven bins at 5×40.8 , sharpest in 0.05–0.08, and between 8 and 26 mrad at 18×137.5 .

Coherent $e^+Li \rightarrow e^+X^6Li(g.s.)$ at the reconstructed level (6R), $e(5) \times 6Li(40.8/u)$, $P_{ZZ} = 0.60$: $\sigma_p = 36 / 380 \mu\text{rad}$ (h / v), rectangle cutout $|p_{x1}| < 0.09$, $|p_{y1}| < 0.93 \text{ GeV}$ [tagging optics]
 $N_{\text{tag}} = 2.3e+06$ (1 yr) vs $5.8e+06$ with the constant 0.20 GeV cut at L_{HA} ; deformation $c_2 = 2a_2$ (Eq. 9 convention), $a_e = 0.010$
 $u_1 = 0.05$, $u_2 = 0.02$ (ZEUS LPS: u_1 its central A_{LT} , u_2 within one error of zero); two-fill ratio fit with the MC template basis; statistical errors only



Money plot 6R — the coherent channel at the reconstructed level (5×40.8 at the tagging optics). **(a)** Raw spin-sorted yields versus the reconstructed recoil azimuth β in the lowest $|t|$ bin, 0.017–0.028 GeV^2 , eight of whose twenty-four β bins the cutout leaves empty: the cutout shape dominates both states identically. **(b)** The modulations $T(\beta)$ and $T(\alpha)$ from the spin-state ratio, the template-fit prediction per β bin (ticks) and the smooth $a_t \cos 2\beta$ and $a_e \cos 2\alpha$ curves. **(c)** The expected modulation map in (α, β) . **(d)** $a_t(t_{ref})$ and a_e in each of the seven reconstructed $|t|$ bins (1 yr open, 10 yr filled) against the injected curves, whose four lowest points — $t_{ref} = 0.023, 0.030, 0.036$ and 0.046 GeV^2 , the last of them the retired window's own lowest bin — are the extrapolation below the anchor described in §4.5.

6 • Systematic effects and the specifications they imply

Each effect was measured on the chain by holding the response fixed and varying one input with common random numbers (inclusive: `money_cos2phi_reco.py --syst-scan`; coherent: exact expected counts with the assumption changed on the fit side only). The response's own Monte-Carlo floor is 0.06–0.28% on Δ at the four spots.

EFFECT	SIZE	ON THE RESULT
electron energy scale $\pm 1\%$	unknown to the analysis	Δ by 0.3–1.0% at the four spots ($d \ln x/d \ln E' = 2 - y$ with the Σ method)
hadronic energy scale, residual after calibration	+2%	$\hat{A}$ by +0.2 to +1.2% (0.1–0.6% per 1%)
hadronic resolution mis-modelled	generated at 30%, corrected with 25%	Δ by 0.5–1.4%

η -dependent calorimetry; ϵ_{elD} η tilt of 0.05	—	$\leq 0.1\%$; 0.02%
bin-centering model dependence	a wrong prior	2.5-8.6% at the spots with an assumed K, $\leq 1.8\%$ (model B) / $\leq 4.6\%$ (toy) with the shape fitted through the response
acceptance harmonic differing between the spin states	$\delta\epsilon_2 = 10^{-3}$	fakes $\hat{A} = 5.6 \times 10^{-4}$ — 6-13% of the projected amplitudes, 2-6 one-year errors
relative luminosity between the states	10^{-3} on the ratio	0.03% (1/3 per unit ratio error, a constant in R)
P_{zz} scale	$\delta P_{zz}/P_{zz}$	1:1 on Δ , exactly (the ratio carries one power of the assumed scale against σ_{p^2} 's two); a pure normalisation — every shape claim is immune
coherent: cutout differing between the states	horizontal half-width differing by 10^{-3} (10^{-2}) between the states	a_t by $-9.1 / -3.2 / -1.8 / -1.1 / -0.5 / -0.5 / -0.5\%$ ($-83.3 / -30.1 / -18.0 / -9.8 / -4.9 / -4.5 / -3.5\%$) in the seven $ t $ bins at 5×40.8 , worst -31.2% (-298.6%) at 10×99.5 and -10.1% (-99.4%) at 18×137.5 , always in the lowest bin, a_e untouched — at the tagging cutout, whose horizontal edge sits inside a shallow part of the spectrum, a $\sim 10^{-3}$ stability of the envelope buys 0.2 to 0.8% in the three bins above 0.08 GeV^2 and 0.9 to 1.9% in 0.05-0.08, but 1.8 to 31.2% in the three bins below 0.05 GeV^2 , where that edge cuts the spectrum steeply (it was +19% in every bin under the earlier slot geometry)
coherent: u_2 assumed wrong	$\delta u_2 = 0.024$ (a ZEUS 1σ)	a_e by $+0.0003$ to $+0.0009$ (3-9%, at most 0.3 one-year errors) over the seven $ t $ bins at 5×40.8 and $+0.0004$ to $+0.0013$ (4-13%) at 18×137.5 , a_t by less than its printed resolution; a wrong u_1 moves only the nuisance a_m — first order through $a_t \delta u_2 (\cos 2\beta)$, so the added low bins are the least affected, a_t being smallest there. Measured in situ instead (next row) this falls to 0.2-6.3% of a_e
coherent: u_1, u_2 measured in situ	fit to the spin-averaged counts against the acceptance MC	$\delta u_2 = 0.0030 / 0.0031 / 0.0035 / 0.0028 / 0.0041 / 0.0071 / 0.0138$ per $ t $ bin at 5×40.8 in one year ($0.0016 / 0.0017 / 0.0020 / 0.0016 / 0.0026 / 0.0049 / 0.0111$ at 18×137.5), 1.7-15 $\times$ tighter than the ZEUS band, the sparsest bin of either configuration the least improved; propagating into a_e as 0.00002 to 0.00063 across the fourteen bins, negligible in quadrature against the 0.00134-0.01427 statistical error. The assumption moves rather than disappears: with the per-bin acceptance free, u is not identifiable at all (u_b multiplies ϵ_b identically), so an in-situ u necessarily rests on the acceptance shape the simulation supplies
coherent: relative luminosity	10^{-3}	$\leq 0.2\%$ on a_t in all seven $ t $ bins — the printed resolution of a_t in the lowest ones — and nothing on a_e

Table 6 — the systematic effects measured on the chain. Not modelled, and outside every error bar here: radiative corrections — §7 bounds the collinear-ISR migration at $+0.5$ to $+1.2\%$ of Δ on the stand-in chain ($+0.4$ to $+0.9\%$ through PYTHIA) in the published generator window and $\leq 2.9\%$ with the low- Q^2 feed-in opened up, but the tensor sector is uncalculated — backgrounds, and the F_2/R model (§7).

Two of the resulting specifications are this programme's to set, because no EIC document addresses a tensor-polarized beam: the tensor-polarization scale, $\delta P_{zz}/P_{zz} \leq 5\%$ (3% has been reached on stored deuterons at the Nuclotron and on the HERMES storage cell, 8-10% is where solid-target work sits; the precedent is collected in plans/00 run 10), propagating 1:1 into the magnitude of Δ and not at all into its shape; and the relative luminosity between the tensor states, $\leq 10^{-3}$, which costs 0.03% and disappears with bunch-by-bunch alternation. The third is the acceptance-harmonic stability of 10^{-4} between the states, or the alternation that makes it exact. The others are asks with owners: $P_{zz} \geq 0.8$ from the source; a hadronic-method $\delta y/y \lesssim 0.3$ at $y \approx 0.01$ from ePIC full simulation, which turns on the calorimeter noise floor at $\Sigma_h = 0.2$ – 0.5 GeV ; the lithium tagging optics and pots that follow it from C-AD and the far-forward group; for the $\alpha \leftrightarrow {}^6\text{Li}$ separation, EICROC's input charge dynamic range and the LGAD gain-suppression curve at $\approx 9 \text{ MIP}$, with a Geant4 study

through the four planes (open question #19, Report 3 Table 9 item 5), against which the four-plane fake rate is $2.3\text{--}3.1 \times 10^{-5}$ at 95% ${}^6\text{Li}$ efficiency (Report 4 §4.1); and a far-forward reconstruction that knows what lithium is, which is what this development supplies.

7 • What the chain does not model

Radiative corrections are not generated in the published response, and for the tensor sector they have never been calculated at all; what collinear initial-state radiation would cost is measured in the paragraph below, and it is the only part of the omission that can be bounded. Backgrounds are absent: no electron-identification purity or pair-symmetric background inclusively, no $\alpha + d$ breakup, vetoes or Z-identification in the coherent sample, and no beam-gas or pile-up. The coherent pseudo-experiment has no electron side and no x_p or Q^2 binning, and its template assumes the linear t shape it fits; the low-count bias of the bin-wise ratio estimator below a few hundred counts per (α, β) cell — already -2.5% at 184 counts per cell and -34.4% at 45, over two hundred draws (Table 5) — is no longer among these — the acceptance-profiled likelihood of §4.3 removes it (Table 5) and is available as an option, with the ratio kept as the default so that the numbers here are the ones the record carries. The $|t|$ window is now a choice the chain makes rather than one it inherits: the seven bins of §4.5 hold 82, 96 and 85% of the tagged sample at the three configurations against 28, 38 and 27% for the $0.05\text{--}0.25 \text{ GeV}^2$ window published until then, and take the combined one-year δa_e from 0.00207 to 0.00121, from 0.00178 to 0.00111 and from 0.00132 to 0.00074. What is still a residual is the one bin below them, $0.006\text{--}0.017 \text{ GeV}^2$, and it stays out on three measurements rather than on taste: the cutout leaves only sixteen of its twenty-four β bins populated at 5×40.8 , eight at 10×99.5 and twelve at 18×137.5 ; its weighted design has condition number 12.82, 47.26 and 14.79, against the range §4.5 quotes for the published bins; and a 10^{-3} fill-to-fill split of the envelope moves its a_t by -43 , -273 and -56% , where no published bin moves by more than 31.2% (Table 6). It would buy 0.00109, 0.00108 and 0.00068, which is not what those three numbers cost. The bins that were added are themselves resolution-dominated — the vertical divergence is 93 MeV and the $0.17\text{--}0.25$ reco bin already has $t_{\text{ref}} = 0.09 \text{ GeV}^2$ — so they buy statistics for the flat a_e far more than shape information for $a_t(t)$, and they are compared against an extrapolation of the anchor rather than against it — as the retired window's lowest bin, at $t_{\text{ref}} = 0.046 \text{ GeV}^2$, already was (§4.5). The structure functions are placeholders behind grid and R1998 backends. The track angular resolution and the calorimeter noise floor are stand-ins that the ePIC full simulation would replace, and the pot aperture is a measurement in one snapshot of the ePIC geometry (`epic-main` at git 9aaa2969, re-measured 2026-08-28) whose 5×41 row stands on the baseline proton lattice for that ring setting, a $\times 0.64$ systematic if a ${}^6\text{Li}$ fill runs the $Z/A = 0.5$ file (Table 2). The tagging optics is a proposal whose own transport has not been computed: its envelope is priced on the measured baseline levers of each configuration, taken to survive the de-squeeze, and a de-squeezed lattice's R_{12} is among the inputs Report 1 §6.1 requests. Each of these is listed in Table 2 or §6 with its leverage, and none changes the closure of §5, which is a statement about the chain given its inputs.

Collinear initial-state radiation is bounded rather than assumed. Every pseudo-event is given a photon of energy fraction z from the exponentiated leading-log electron structure function $D(z, Q^2) = (t/2)z^{t/2-1}S(t) - (t/4)(2-z)$, $t = (2\alpha/\pi)[\ln(Q^2/m_e^2) - 1]$, $S(t) = e^{(t/2)(3/4-\gamma_E)}/\Gamma(1+t/2)$ [21,22]; the hard vertex then sits at $\hat{s} = (1-z)s$ while the analysis reconstructs with the nominal beam. The truncated exponentiation normalises to $1 + O(t^2)$, the residual being 7×10^{-4} at the mid configuration's rate-weighted $\langle t \rangle = 0.070$, its mean Q^2 being 4.4 GeV^2 , where 27% of the selected rate radiates above $z = 10^{-4}$, $\langle z \rangle = 0.025$ and $\langle z \rangle = 0.092$ among those that do. Three of the four effects this could have are exactly zero. The covariant azimuth is invariant under $k \rightarrow (1-z)k$ for a massless target — $\cos \varphi'$ and $\sin \varphi'$ both acquire the factor $[2ac((1-z)a-b)]^{-1/2}$ with $a = P \cdot k$, $b = P \cdot k'$, $c = k \cdot k'$, and the arctangent divides it out — so radiation fakes neither a $\cos \varphi'$ nor a $\cos 2\varphi'$ modulation. The residual is 3.6×10^{-15} rad — double-precision roundoff — over a 2×10^4 -event sample flat in $(\log x, \log y)$ with z up to 0.9; over the 1.84×10^6 events of the response itself, where the physical ${}^6\text{Li}$ mass leaves the $O(\gamma^2)$ term of the projector, it reaches 2.6×10^{-2} rad and dilutes $\cos 2\varphi'$ by 9×10^{-8} rate-weighted. An unpolarized-lepton QED correction is common to the two fills by construction and cancels in the spin-state ratio of §4.3. And x is exact: the nominal s the analysis divides by exceeds the true $\hat{s}$ by $1/(1-z)$, the electron-method Q^2 exceeds the hard Q^2 by the same $1/(1-z)$, and y_Σ knows no beam energy [6], so $x = Q_e^2/(s y_\Sigma)$ is the hard x identically. What is left is the Q_e^2 label, high by $1/(1-z)$, and the migration it

causes. Measured with common random numbers — the radiation draws from its own stream, so the two responses sit on the same pseudo-events — at 1600 pseudo-events per cell, the four sweet spots lose 0.015 to 0.044 in purity (0.653 → 0.638, 0.633 → 0.613, 0.679 → 0.659, 0.684 → 0.640) and 0.005 to 0.019 in efficiency, and Δ corrected with an ISR-free bin-centering is biased by $+0.62 \pm 0.03$, $+0.50 \pm 0.02$, $+0.94 \pm 0.03$ and $+1.22 \pm 0.02\%$. A single response draw scatters by 4 to 14% of the bound — the seed-to-seed standard deviation is 0.05 to 0.09 percentage points and one seed can sit two of those from the mean — so those are means over eight response seeds and the errors are the standard errors of that mean, not a Monte Carlo floor. The generator window truncates the feed-in, since an event whose Q^2 lies below it cannot radiate into an analysis bin; opening it moves the worst spot to 1.87, 2.81, 2.26 and 2.34% at $Q^2 = 0.35, 0.15, 0.05$ and 0.02 GeV^2 , so the bound saturates in a 1.8–2.8% band rather than at a single value, and $\leq 2.9\%$ is what the $\leq 5\%$ gate of plans/07 WP4 is read against. It passes by a factor 1.7. The number is not an artefact of the Gaussian y stand-in: run through the PYTHIA hadronic final state with the calibrated scale it is $+0.38, +0.44, +0.46$ and $+0.88\%$. The method comparison behind the choice of estimator is sharp, and has to be read at a stated y because the electron-method ratios go as $(y+z)(1-z)/y$. At $z = 0.092$ and the rate-weighted $\langle y \rangle = 0.189$ of the whole selected sample the observed-to-hard ratios (Q^2, y, x) are (1.102, 1.351, 0.740) for the electron method, (1.000, 1.000, 0.908) for Σ , (0.976, 0.908, 0.976) for Jacquet–Blondel, (1.214, 1.000, 1.102) for the double angle and (1.102, 1.000, 1.000) for the mixed pair the chain uses. At the sweet spots themselves, where $y = 0.0101, 0.0254, 0.0111$ and 0.0254 , the same photon multiplies the electron-method y by 9.2, 4.2, 8.5 and 4.2 and collapses x_e to 0.11–0.24 of the hard x , quite apart from resolution; every other row is unchanged except Jacquet–Blondel's Q^2 , which rises to 0.998. Only the mixed row is 1.000 at every y . $Q^2_\Sigma = p_{T,e}^2/(1-y_\Sigma)$ uses no beam energy either, so labelling the bin with it instead of Q_e^2 would remove the migration entirely at the price of the Σ resolution on the label; that is a change to the chain and is not made here. Neither is the $E - p_z$ window every HERA analysis used against radiative events: the visible sum is $2(1-z)E_e$ and is already reconstructed, as $\Sigma_h + E'(1 - \cos \theta) = E'(1 - \cos \theta)/(1 - y_\Sigma)$, and requiring it within 15% of $2E_e$ keeps 87% of the non-radiative selected rate on the 25% Gaussian y stand-in and 98% through the PYTHIA final state with the calibrated Σ scale (82% and 93% of the radiative rate), the loss falling almost entirely above $y = 0.2$, where the two chains put 99.5% and 97.2% of the discarded non-radiative rate — the retention below it is 1.000 and 0.998 on either chain, against 0.78 and 0.27 above it on the stand-in and 0.97 and 0.88 on PYTHIA — and amounting to 0.06% or less at the four sweet spots themselves; the mitigated bias is $+0.23, +0.16, +0.22$ and $+0.18\%$ on the stand-in chain and $+0.17, +0.19, +0.20$ and $+0.16\%$ on the PYTHIA one, independent of the generator window. The chain does not apply the window: the gate already passes at $\leq 2.9\%$ without it, it is free only where these bins sit, and the rate it removes is the $y > 0.2$ rate that carries the low- x end of every Q^2 slice. Outside all of this and unbounded: the tensor-sector radiative correction, which has never been calculated (plans/05 §5.5) and which no unpolarized study can stand in for; the polarized-lepton correction, irrelevant to an unpolarized beam; wide-angle real emission, which carries no logarithm and does not factorise into a beam-energy shift; final-state radiation, largely reabsorbed by the calorimeter cluster; and the elastic and quasi-elastic radiative tails, which the $W^2 \geq 10 \text{ GeV}^2$ cut removes.

References

- [1] R. L. Jaffe, A. Manohar, *Nuclear gluonometry*, Phys. Lett. B 223 (1989) 218.
- [2] P. Hoodbhoy, R. L. Jaffe, A. Manohar, *Novel effects in deep inelastic scattering from spin-one hadrons*, Nucl. Phys. B 312 (1989) 571.
- [3] A. Bacchetta, M. Diehl, K. Goeke, A. Metz, P. J. Mulders, M. Schlegel, *Semi-inclusive deep inelastic scattering at small transverse momentum*, JHEP 02 (2007) 093, hep-ph/0611265 (Eqs. 2.4–2.5).
- [4] A. Bacchetta, U. D'Alesio, M. Diehl, C. A. Miller, *Single-spin asymmetries: the Trento conventions*, Phys. Rev. D 70 (2004) 117504; M. Diehl, S. Sapeta, Eur. Phys. J. C 41 (2005) 515.
- [5] F. Jacquet, A. Blondel, in *Proceedings of the Study of an ep Facility for Europe*, DESY 79/48 (1979) 391.
- [6] U. Bassler, G. Bernardi, *On the kinematic reconstruction of deep inelastic scattering at HERA: the Σ method*, Nucl. Instrum. Meth. A 361 (1995) 197.
- [7] R. Abdul Khalek et al., *Science requirements and detector concepts for the Electron-Ion Collider: EIC Yellow Report*, Nucl. Phys. A 1026 (2022) 122447.
- [8] A. Jentsch, Z. Tu, C. Weiss, Phys. Rev. C 104 (2021) 065205, arXiv:2108.08314 (far-forward acceptance model).
- [9] H. Mäntysaari, F. Salazar, B. Schenke, C. Shen, W. Zhao, *Spatial imaging of polarized deuterons at the Electron-Ion Collider*, Phys. Lett. B 858 (2024) 139053, arXiv:2408.13213.

- [10] W. Chang, E.-C. Aschenauer, A. Jentsch, A. Kumar, Z. Tu, Z. Yin, Phys. Rev. D 113 (2026) 032018, arXiv:2511.05638 (IR-8 secondary focus).
- [11] This repository: Report 1 (`reports/cos2phi_money_plots_report`), Report 3 (`reports/eic_epic_reference`), `evgen/polligen/{reco,hfs,recopseudo}.py` , `docs/reproduction_manual.md` , `docs/code_review_2026-08-28.md` .
- [12] W. Cosyn, B. Roldan Tomei, A. Sosa, A. Zec, *Polarization options in inclusive DIS off tensor polarized deuteron*, Eur. Phys. J. A 61 (2025) 83, arXiv:2410.12764.
- [13] ZEUS Collaboration, S. Chekanov et al., Nucl. Phys. B 816 (2009) 1, arXiv:0812.2003 (leading-proton azimuthal asymmetries; beam-spread-dominated LPS resolution).
- [14] N. N. Nikolaev, A. V. Pronyaev, B. G. Zakharov, arXiv:hep-ph/9812212 (the y dependence of the diffractive azimuthal terms).
- [15] A. Jentsch (ePIC), *Far-Forward Detectors and Physics with ePIC @ the EIC*, DIS 2023 (beam effects dominate the far-forward momentum smearing).
- [16] ATHENA Collaboration, JINST 17 (2022) P10019, arXiv:2210.09048 (Sec. 3.1, Fig. 22: hadronic-method resolution at low y ; Table 5: electron identification).
- [17] S. Maple (ePIC), *Tracking and inclusive DIS reconstruction with the ePIC detector at the EIC*, seminar, University of Birmingham, 11 December 2024 (the 25% hadronic smearing of the kinematic-fit study).
- [18] M. Arratia, D. Britzger, O. Long, B. Nachman, Nucl. Instrum. Meth. A 1025 (2022) 166164, arXiv:2110.05505 (Table 1: the classical methods; Sec. 5: noise and acceptance at low y).
- [19] K. Abe et al. (E143), Phys. Lett. B 452 (1999) 194, arXiv:hep-ex/9808028 (the R1998 world fit).
- [20] T. Sjöstrand et al., *An introduction to PYTHIA 8.2*, Comput. Phys. Commun. 191 (2015) 159; PYTHIA 8.3 (2022), the version used is 8.311.
- [21] E. A. Kuraev, V. S. Fadin, *On radiative corrections to e^+e^- single photon annihilation at high energy*, Sov. J. Nucl. Phys. 41 (1985) 466 (Yad. Fiz. 41 (1985) 733) — the leading-log electron structure function with the soft series exponentiated.
- [22] O. Nicrosini, L. Trentadue, Phys. Lett. B 196 (1987) 551; M. Skrzypek, Acta Phys. Polon. B 23 (1992) 135 (Eq. 2: the $D(z)$ used here).

Appendix A · Revision history and the 2026-08-28 review

DATE REVISION

2026-09-06 Sourcing pass. Table 2's tracking row names the Yellow Report sources that actually differ backward — Table 11.2 against Fig. 8.3 / Table 8.20 — as Report 3 Table 8 now does; Fig. 8.3 was read as an image for the first time on 2026-09-06; §4.1 records the $\hat{\kappa}$ subtraction of the tensor leakage, default off.

2026-09-02 **Consistency revision** from a cross-document review of the five reports against the code, the plans and the manual, adjudicated finding by finding in `docs/consistency_review_2026-09-02.md` (its identifiers in parentheses). §2 states the b_1 , b_2 sign convention at the master formula itself — the simulation has carried `TENSOR_LL_SIGN` = -1 since 2026-08-29, the convention every published b_1 is quoted in — and says that the tensor polarization weights the inclusive rate as well as its azimuth, the hadronic final state being blind to it (F002, F013, F200). Table 2: the Yellow Report hadron table is its strong-hadron-cooling one, at the start-of-store emittance (C001); the ${}^6\text{Li}$ menu cites this repository rather than the ZEUS leading-proton paper (F082); the spin axis is the arcs' stable direction assumed delivered at IP6 (F137); η = -2.9 to -1.6 is the mid configuration's four spots, against -3.4 to -0.9 over the twelve (F134); the angular placeholder is given per spot (F201) and ϵ_{elD} 's single anchor and low forward tail are stated (F204); the nominal forward reach recovers a third of the loss, $0.170 \rightarrow 0.115$ at the first spot (F199); the recoil's own dispersion is stated rather than modelled, what the tagging optics folds in being the beam's momentum spread (F083); the fixed-pot tag reads 10^{-5} – 10^{-19} against 9×10^{-10} – 6×10^{-27} at the published optics (F012); u_1 and u_2 are placed against the ZEUS central values rather than at a 1σ edge (C003); the R1998 shift at the second spot is $+18\%$ (F190); and the Z-identification ask becomes EICROC's input charge range and the gain-suppression curve, with the $\alpha + d$ partner veto as a second lever (F202). §4.1 conditions the $O(\gamma^2)$ tensor leakage on the $b_3 = b_4 = 0$ the kernel defaults to (-0.052% at a tenth of b_2) and records the $A_{\parallel}$ target-mass term as the generator default since 2026-08-29 (F174, F006). §4.2 and Table 3: all four sweet spots sit at $y = 0.010$ – 0.025 , not three (F062, F089); $\delta Q^2/Q^2$ is given as an rms through the response beside the bin-centre formula (F198); $E' = 9.93 / 9.77 / 9.97$ GeV and $\theta' = 108$ mrad at the second spot (F207, F208); the figure pointers follow the 2026-08-29 renumbering (F090). §4.4 restates the unfold scan from the run that produces it — 20–29% with an assumed prior, folded residuals $4.0 / 5.7 / 5.8\%$ (model B) and $4.0 / 4.3 / 3.6\%$ (toy), the spot quadruples ending $+3.1$ and $+2.9$ — and Table 6's row with it, 2.5–8.6% (F069, F132, F197). §4.5 and Figure 3: the fixed-aperture tag is sourced to Report 3 §5 and the caption accounts for the markers panel (c)'s floor hides (F131); the tagged fraction 37.6% is per recoil above t_{min} , against Table 5's 36% per recoil produced (F210); $\langle \cos 2\theta \rangle = -0.49$ and -0.95 are in the true recoil azimuth, the divergence smearing diluting the first to the -0.31 money plot 6R annotates (F087). §5: Table 4's third row carries Report 1's current 2.7×10^{-4} (F064); the reconstructed-to-truth error ratio is 0.60–0.72 over the two hadronic- y models and 0.60–0.67 with the Gaussian stand-in, the reconstruction efficiency being read before the analysis cuts (F066, F213, F165); the prior spread $3.7 / 2.1 / 1.3\%$ is the ten-year decomposition, against $3.9 / 2.0 / 1.4\%$ in one year (F196); the abstract's "unbiased above 10^5 recoils" becomes a 1.5% bound and Table 5's caption is scoped to its twenty draws (F170); Figure 4's panel (c) purities are migration-only and its panel (d) count truth-level (F205, F206). §6 points at plans/00 run 10 for the polarimetry precedents rather than at Report 1 §7 (F067); §7 states the pot aperture as a 2026-08-28

measurement in the current geometry, lattice-conditional at 5×41 (F005), the sweet-spot y values as $0.0101 / 0.0254 / 0.0111 / 0.0254$ (F033), and $\langle t \rangle = 0.070$ and $\langle Q^2 \rangle = 4.4 \text{ GeV}^2$ as separate averages (F195). The rows below are corrected where they misstate what they record: the fourth sweet spot named by its bin (F081), the target-mass default (F006), and the cutout floor separated into $3.8\text{--}6.3\times$ below the retired 0.05 GeV^2 edge and $1.3\text{--}2.2\times$ below the adopted one (F194). The figures are regenerated on the same libraries with corrected labels — the hadronic-resolution legend names the earlier chain rather than the published 5R/7R, money plot 7R's title fits its panel, and 6R's u_1, u_2 curves are no longer called 1σ bounds (F209, F211, C003); no number moves with them. Wording suggestions of the same review applied: F030, F032, F063, F065, F084, F088, F091 — the open rank-2 question on Δ 's dilution, the alignment axis referenced to the beam at the master formula, the ISR bound labelled per hadronic- y chain, the tagging optics priced on each configuration's measured baseline levers, Table 5's unbiased bins scoped to 5×40.8 , the $O(\gamma^2)$ leakage stated per closure configuration, and the EMCal resolution's factor-2 optimism at the fourth sweet spot.

2026-08-29 Verification pass over the round. §4.1's $O(\gamma^2)$ azimuth bound restated on a re-run of `target_mass_bound.py` : 1.4, 1.2 and 5.9 mrad at the mid, top and low configurations, against the 1.3 and 6.0 the text had rounded to. Table 3's fourth sweet spot, $(0.141, 14.3 \text{ GeV}^2)$, corrected to the run that produced it — $\Sigma, \text{JB}, \text{DA} = 0.15, 0.15, 0.21$ and $\delta y/y(e' \text{ alone}) = 0.46$, and 1.07 at the third spot. The acceptance-harmonic fake is five to eleven times the projected amplitudes, as Report 1 states it. Table 5's caption no longer refers to sine columns the table does not carry. §4.5 notes that the fourth extrapolated bin is marginal at 18×137.5 , $t_{\text{ref}} = 0.0498$ against the anchor's lowest point at 0.050. The $Z/A = 0.5$ lattice edge is quoted as the applied 1.61 mrad. §2's six-stage list is set as prose. The figures are renumbered into order of appearance — the two hadronic figures 1 and 2, the coherent figure 3, the inclusive figure 4 — and references [8], [10], [11], [16], [17], [19] and [20], defined but uncited, are cited at their anchors. §4.5's cutout harmonics are restated from the run that produces them: the fake $(\cos 2\beta)$ at the measured aperture is -0.95 , the value the 2026-08-28 re-measurement implies and the one its two sentence-mates, 9×10^{-10} and 7×10^{-14} , were already carrying, and at the tagging cutout it is -0.49 on the per-configuration levers, rather than the -0.27 the retired $72.7 \mu\text{rad}$ envelope had left behind; the tagging-optics tag there is 37%, and Table 5's 6R row reads 36% at the reconstructed level. The 7R stand-in error at $Q^2 = 1.14 \text{ GeV}^2$ is 2.4×10^{-3} and the PYTHIA one at 3.13 GeV^2 1.2×10^{-3} , and Table 4's spot-1 PYTHIA $\bar{\alpha} \pm 0.12$, each a rounding of the printed value rather than a change in it. The abstract's tagged-recoil range follows Table 5's own row headers, $2.3\text{--}6.0 \times 10^6$ a year, the lower end having moved with the per-configuration transport adopted the same day. §4.1's $b_1\text{--}b_4$ leakage into $\cos 2\phi'$ is stated as a measurement rather than a bound from the same day: the exact finite- γ tensor kernel makes it -0.033% of the amplitude at this configuration's sweet spots, with the sign that cancels, in place of the $\leq 0.15\%$ the bound claimed — first published that day as -0.23% , and corrected within it once the two finite- γ rows of the source's own Table 1 were made the kernel's anchor and fixed the twist-3 channel's sign as the cancelling one; and the sign of the by-product κ in §4.3 is fixed by the tensor convention settled with it. The finite- γ longitudinal kernel implemented default-off the day before is the generator's default from this date (`InclusiveKernel(target_mass=True)`), the author decision of 2026-08-29; `target_mass=False` remains the massless limit, bit-for-bit), and §4.1 says so beside the tensor switch, whose default is the opposite; no number of this report moves, since no Δ here is built from $A_{||}$.

2026-08-28 Table 2's luminosity row separates two shares that the word had covered indistinguishably: the $0.5/0.5$ division between this measurement's own spin states, which divides whatever luminosity it is given and which the analysis knows only to 10^{-3} , and the programme share of plans/07 WP2, which divides the year between observables, far-forward optics and isotopes; the two multiply, and neither is chosen here. Both closures below are run at the full $10 \text{ fb}^{-1}/\text{u}$.

2026-08-28 The pot-aperture input row re-measured in the current ePIC geometry (`epic-main` at git 9aaa2969, 0.05 mrad ladders at four azimuths; plans/09 B1): $2.50/\text{shut}$, $1.51/2.12$ and $0.53/0.92 \text{ mrad}$ (h/v) against the September-2024 $2.00/3.00$, $1.35/3.00$ and $1.03/2.30$, so the tag with the pots parked at that edge runs $10^{-5}\text{--}10^{-19}$ where the retired geometry gave $10^{-6}\text{--}10^{-17}$. One number downstream of it moves, §4.5's tagged fraction at the Yellow Report optics, which was 7×10^{-8} , 6×10^{-27} and 2×10^{-17} on that geometry; the closures do not, being run at the tagging optics with the pots following the envelope, where the aperture does not enter.

2026-08-28 **The coherent $|t|$ window widened to the seven bins $0.017\text{--}0.25 \text{ GeV}^2$** (plans/08 item 4, whose last surviving lever this was: adopted, not superseded). The $0.05\text{--}0.25 \text{ GeV}^2$ window the earlier revisions published was inherited from the pre-measurement slot geometry and threw away 62–73% of the tagged sample; its lower edge was never the apparatus, the cutout's own floor lying a factor 3.8 to 6.3 below it, and a factor 1.3 to 2.2 below the 0.017 GeV^2 edge adopted here (§4.5). On the wider window the combined one-year δa_e falls from $0.00207 / 0.00178 / 0.00132$ to $0.00121 / 0.00111 / 0.00074$ at the three configurations, and Tables 5 and 6 and §7 are re-measured on it throughout — the estimator's low-count bias, the envelope-split sensitivity, the in-situ δu_2 and the pot-roll resolution all now read per seven bins, the added ones being the sparsest in β and the most exposed to the cutout edge. The bin below the window, $0.006\text{--}0.017 \text{ GeV}^2$, is excluded on three measurements rather than on taste (§7), and the retired window survives as `--t-edges 0.05,0.08,0.12,0.17,0.25` in both coherent scripts under its own output stem, so every number the record carries from before this revision is reproducible. Each 6R run now prints a per-bin `design:` line — counts per (α, β) cell, β bins populated, condition number, and the rank `reco._harmonic_rank_guard` measures on the weighted design, which both fitters now return — which is where the conditioning statements of §4.5 and §7 are read off. The anchor caveat is stated per bin as four of the seven and not three, since a_t is quoted at t_{ref} and the $0.05\text{--}0.08$ bin's is 0.046 GeV^2 . Two tests added; the published figure stems do not move.

2026-08-28	§4.1's $O(\gamma^2)$ bound on the lab-angle azimuth shortcut measured rather than asserted: at most 1.4 and 1.2 mrad at the sweet spots of the mid and top configurations and 5.9 mrad at those of the low one, against $\gamma^2 \leq 0.0057, 0.0049$ and 0.025 there, replacing a single $\gamma^2 \leq 0.02$ quoted for all three and an unmeasured 5 / 12 mrad pair. The finite-γ longitudinal kernel $A_{\parallel} = D_{\gamma}(A_1 + \eta A_2)$ that this γ^2 belongs to is now implemented behind <code>InclusiveKernel(target_mass=True)</code>, default off, and bounded on every observable built from $A_{\parallel}$ by <code>evgen/scripts/target_mass_bound.py</code>. No published number moves.
2026-08-28	Housekeeping from the integrated review: the reference list, which had skipped [15] and [16] since its first numbering, is renumbered so that the citations run 1–22 without a gap; §7 states the onset of the ratio's low-count bias at a few hundred counts per (α, β) cell rather than ~ 30, which is what Table 5 measures; Table 5's caption names the default 12×24 grid that the coarser binnings are compared against; and <code>harmonic_ratio_fit_2d</code> now returns <code>err_const</code>, so the ratio and the likelihood share the return shape the docstring advertises.
2026-08-28	The coherent estimator's low-count bias closed, and the unpolarized harmonics measured in situ (plans/08 A12, A3). The inversion of the bin-wise spin-state ratio is curved — concave for the flip plan, whose $\bar{P} = -0.3$ sets the sign — over the range the ratio explores at low counts, which put a negative per-bin offset of order $1/\nu$ on $\hat{T}$ and — the acceptance being flat in α and modulated $\times 25$ in β — landed it on a_t and the constant. The acceptance-profiled Poisson likelihood added alongside it (§4.3) is exactly the conditional multinomial given the bin totals, hence unbiased at any occupancy; Table 5 gains its column and §7 loses the corresponding caveat. The ratio remains the default, so no number previously published moved: the two agree to machine precision on exact counts and to 0.24% on the Asimov errors. (u_1, u_2) are now also fittable from the spin-averaged counts against the acceptance MC, which replaces the ZEUS 1σ band in Table 6 by a 1.7–15$\times$ tighter in-situ constraint and makes explicit that an in-situ u rests on the acceptance shape rather than on the data alone. The two-azimuth basis under a strongly anisotropic acceptance, listed as open since the aperture measurement, was closed by measurement: the design is near-orthogonal at the tagging optics (condition number 1.9–3.4 over the four bins published then, 1.85–10.87 over the seven published now, $\text{corr}(a_e, a_t) < 0.005$) and the anisotropy regime the item was written for became an empty sample when the ion energies were γ-matched.
2026-08-28	Radiative corrections bounded rather than assumed (plans/07 WP4, plans/08 D3, plans/02 step 1.4): <code>polligen/radiative.py</code> with the exponentiated leading-log photon spectrum, a default-off <code>RecoResponse(isr=...)</code> migration hook and <code>money_cos2phi_reco.py --isr</code>. §7 carries the result — collinear ISR fakes no $\cos \phi'$ or $\cos 2\phi'$, cancels in the spin-state ratio, leaves x exact and moves only the Q^2_e label — with the gate verdict (Δ biased by $+0.5$ to $+1.2\%$ in the published generator window, $\leq 2.9\%$ once that window is opened below $Q^2 = 0.15 \text{ GeV}^2$ and the low-Q^2 feed-in saturates, $\leq 0.25\%$ behind an $E - p_z$ window, against a 5% gate) and the method comparison; a new row in Table 2. The bound is quoted as a mean over eight response seeds, one draw of the response scattering by about its own size; the method comparison is quoted with the y it is evaluated at, the electron-method rows being a strong function of it; and the azimuth residual is quoted per sample, the flat test sample and the response differing by two orders. 25 tests added; no published number moves, the hook being off by default. The $E - p_z$ retention published with that revision (87% / 82%) was the 25% Gaussian y stand-in's, quoted without saying so; it is now given per chain, per analysis bin and in bands of y, the PYTHIA chain with the calibrated Σ scale keeping 98% / 93%, and the window is stated as a contingency rather than a default (five further tests, 30 in all).
2026-08-28	Rewritten after a full review of the toolchain (nine independent readings of the code against the questions: is each step right, does any analysis step read information an experiment cannot have, is every assumption stated; findings adversarially verified and recorded in <code>docs/code_review_2026-08-28.md</code>). What was wrong and is fixed: the hadronic-scale calibration was looked up at each pseudo-event's true (x, Q^2) cell — now a map in reconstructed $(x_{\text{mixed}}, Q^2_e)$ bins built from the library's own reconstructed events, purities 0.56 / 0.59 / 0.64 / 0.75 against 0.52 / 0.59 / 0.59 / 0.76 before; the PYTHIA library carried PYTHIA's default $\hat{m} \geq 4 \text{ GeV}$, which removed $x < 16/s$ (39% of the selected rate) — regenerated with $\hat{m} \geq 0.5 \text{ GeV}$; the estimator's bin variance vanished for a bin populated by one spin state (an infinite weight, and the "rank-deficient" t bins of the earlier Table 3) — now the expected-count form; the ratio inversion's error Jacobian dropped a second-order term; the hadron acceptance was applied in the head-on frame ($+0.7$–1.5% of Σ_h in the lab); the library's target-mass term was not carried onto the pseudo-events (2% of Σ_h at $y = 0.01$); the $e + p / e + n$ merge was by count rather than cross section; the beam-energy spread had the same sign on Q^2_e and $1 - y_e$; the bins of the 5R panels were selected partly on the truth amplitude; the coherent pseudo-experiment ran on the legacy 73 μrad divergence with a 0.73 mrad "near-beam approach" — now the tagging optics of Report 1 with the Yellow Report divergences, at all three configurations, with a per-fill perturbation that acts on the binding cut; and the "residual template bias" on a_e was a 2σ fluctuation, shown by the ensemble of Table 5. The stand-in numbers of Table 4 did not move; the calibrated ones moved within their errors. Stale statements removed: the pre-correction sweet spots of the old Table 1, coherent numbers at 50 GeV/u, the 73/149 μrad footer, the 25% stand-in as the published hadronic y, the $x \rightarrow x/(1 - z)$ statement about radiation.
2026-08-24 – 27	First version and its revisions: the covariant azimuth and the head-on frame, the electron-method limit and the hadronic methods, the spin-state ratio, the coherent two-azimuth fit with the acceptance-weighted template basis and the sine null columns, the PYTHIA 8 final state through the hadron-side response, the Roman-Pot aperture measured in the ePIC geometry, the forward-folded bin-centering fit, the per-fill acceptance and the assumed nuisances as fit arguments, the γ-matched beam energies, the 1:1 propagation of the P_{zz} scale, and the consolidated

assumptions and specifications. The numbers of those revisions that this one supersedes are recorded in plans/00 and plans/08.

Conventions: head-on frame with the ion along $+z$ and the electron along $-z$; θ' is the scattered-electron angle from the electron beam; per-nucleon (x, y, s) throughout; whole-nucleus four-momenta with the physical nuclear mass where a mass matters. Beam divergences are the Yellow Report Table 10.1 proton values with the γ -matched species step, and the tagging optics of Report 1 §6.1 where stated. Built by `reports/build_report.py`; figures from the `evgen/scripts` named in the text; display mathematics typeset at build time.